



**Image credit:** VantageCircle (<https://blog.vantagecircle.com/employee-attrition/>)

# WELCOME!

Welcome to "**Employee Churn Analysis**" study. In this study, we will be able to build our own classification models for a variety of business settings.

Also we will learn what is Employee Churn?, How it is different from customer churn, Exploratory data analysis and visualization of employee churn dataset using **\*matplotlib\*** and **\*seaborn\***, *model building and evaluation using python* **scikit-learn\*** package.

We will be able to implement classification techniques in Python. Using Scikit-Learn allowing us to successfully make predictions with the Random Forest, Gradient Descent Boosting , KNN and CatBoost algorithms.

At the end of the project, we will have the opportunity to deploy your model using *Streamlit*.

## 1 - DATA

In this project we have HR data of a company. A study is requested from us to predict which employee will churn by using this data.

The HR dataset has 14,999 samples with various information about the employees. In the given dataset, we have two types of employee one who stayed and another who left the company. This given dataset will be used to predict when employees are going to quit by understanding the main drivers of employee churn.

**For a better understanding and more information, please refer to DataCamp (<https://www.datacamp.com/community/tutorials/predicting-employee-churn-python>) and Kaggle Website (<https://www.kaggle.com/c/employee-churn-prediction/data>)**

### 1.1 Context

**"Analyze employee churn. Find out why employees are leaving the company, and learn to predict who will leave the company.." DataCamp (<https://www.datacamp.com/community/tutorials/predicting-employee-churn-python>)**

Employee turn-over (also known as "employee churn") is a costly problem for companies. The true cost of replacing an employee can often be quite large. A study by the Center for American Progress found that companies typically pay about one-fifth of an employee's salary to replace that employee, and the cost can significantly increase if executives or highest-paid employees are to be replaced. In other words, the cost of replacing employees for most employers remains significant. This is due to the amount of time spent to interview and find a replacement, sign-on bonuses, and the loss of productivity for several months while the new employee gets accustomed to the new role.

In the past, most of the focus on the "rates" such as attrition rate and retention rates. HR Managers compute the previous rates try to predict the future rates using data warehousing tools. These rates present the aggregate impact of churn, but this is the half picture. Another approach can be the focus on individual records in addition to aggregate.

There are lots of case studies on customer churn are available. In customer churn, you can predict who and when a customer will stop buying. Employee churn is similar to customer churn. It mainly focuses on the employee rather than the customer. Here, you can predict who, and when an employee will terminate the service. Employee churn is expensive, and incremental improvements will give significant results. It will help us in designing better retention plans and improving employee satisfaction.

### 1.2 About The Features

**We can describe 10 attributes (features) in detail as:**

- **\*satisfaction\_level** :\* It is employee satisfaction point, which ranges from 0-1.
- **\*last\_evaluation** :\* It is evaluated performance by the employer, which also ranges from 0-1.
- **\*number\_projects** :\* How many of projects assigned to an employee?
- **\*average\_monthly\_hours** :\* How many hours in averega an employee worked in a month?
- **\*time\_spent\_company** :\* time\_spent\_company means employee experience. The number of years spent by an employee in the company.
- **\*work\_accident** :\* Whether an employee has had a work accident or not.
- **\*promotion\_last\_5years** :\* Whether an employee has had a promotion in the last 5 years or not.
- **\*Departments** :\* Employee's working department/division.
- **\*Salary** :\* Salary level of the employee such as low, medium and high.
- **\*left** :\* Whether the employee has left the company or not.

## 1.3 What The Problem Is

First of all, to observe the structure of the data, outliers, missing values and features that affect the target variable, we must use exploratory data analysis and data visualization techniques.

Then, we must perform data pre-processing operations such as **\*Scaling\*** and **\*Label Encoding\*** to increase the accuracy score of Gradient Descent Based or Distance-Based algorithms. we are asked to perform **\*Cluster Analysis\*** based on the information you obtain during exploratory data analysis and data visualization processes.

The purpose of clustering analysis is to cluster data with similar characteristics. We are asked to use the **\*K-means\*** algorithm to make cluster analysis. However, you must provide the K-means algorithm with information about the number of clusters it will make predictions. Also, the data we apply to the K-means algorithm must be scaled. In order to find the optimal number of clusters, we are asked to use the **\*Elbow method\***. Briefly, try to predict the set to which individuals are related by using K-means and evaluate the estimation results.

Once the data is ready to be applied to the model, we must **\*split the data into train and test\***. Then build a model to predict whether employees will churn or not. Train our models with our train set, test the success of our model with our test set.

Try to make our predictions by using the algorithms **\*Gradient Boosting Classifier, K Neighbors Classifier, Random Forest Classifier, and CatBoost Classifier. We can use the related modules of the *scikit-learn*\*\*** library. We can use scikit-learn **\*Confusion Metrics\*** module for accuracy calculation. We can use the **\*Yellowbrick\*** module for model selection and visualization.

In the final step, we will deploy your model using Streamlit tool.

## 1.4 Project Structure & Tasks

### 1. Exploratory Data Analysis

- Importing Modules
- Loading Dataset
- Data Insights

### 2. Data Visualization

- Employees Left
- Determine Number of Projects
- Determine Time Spent in Company
- Subplots of Features

### 3. Data Pre-Processing

- Scaling
- Label Encoding

### 4. Cluster Analysis

- Find the optimal number of clusters (k) using the elbow method for for K-means.
- Determine the clusters by using K-Means then Evaluate predicted results.

### 5. Model Building

- Split Data as Train and Test set
- Built Gradient Boosting Classifier, Evaluate Model Performance and Predict Test Data
- Built K Neighbors Classifier and Evaluate Model Performance and Predict Test Data
- Built Random Forest Classifier and Evaluate Model Performance and Predict Test Data

### 6. Model Deployment

- Save and Export the Model as .pkl
- Save and Export Variables as .pkl

## 2 - LIBRARIES NEEDED IN THE STUDY

```
In [1]: # 1-Import Libraies

import pandas_profiling
# import pyforest

import ipywidgets
from ipywidgets import interact

import numpy as np
import pandas as pd

import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
import matplotlib.ticker as mticker
import squarify as sq

# Importing plotly and cufflinks in offline mode
import plotly
import plotly.express as px
import cufflinks as cf
import plotly.graph_objs as go
import plotly.offline as py
from plotly.offline import iplot
from plotly.subplots import make_subplots
import plotly.figure_factory as ff
cf.go_offline()
cf.set_config_file(offline=False, world_readable=True)

# !pip install termcolor
import colorama
from colorama import Fore, Style # makes strings colored
from termcolor import colored
from termcolor import cprint

from wordcloud import WordCloud

import scipy.stats as stats
from scipy.cluster.hierarchy import linkage, dendrogram
import statsmodels.api as sm
import statsmodels.formula.api as smf
import missingno as msno

import datetime as dt
from datetime import datetime

import optuna

from sklearn.cluster import KMeans, AgglomerativeClustering
from sklearn.compose import make_column_transformer, ColumnTransformer
from sklearn.metrics import ConfusionMatrixDisplay
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.dummy import DummyClassifier
from sklearn.ensemble import RandomForestRegressor, RandomForestClassifier, GradientBoostingRegressor
from sklearn.ensemble import ExtraTreesRegressor, AdaBoostClassifier, GradientBoostingClassifier, ExtraTreesClassifier
from sklearn.feature_selection import SelectKBest, SelectPercentile, f_classif, f_regression, mutual_info_regression
from sklearn.impute import SimpleImputer, KNNImputer
from sklearn.linear_model import LinearRegression, Lasso, Ridge, ElasticNet, LogisticRegression
from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error
from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
from sklearn.metrics import make_scorer, precision_score, precision_recall_curve
from sklearn.metrics import roc_auc_score, roc_curve, f1_score, accuracy_score, recall_score
from sklearn.metrics import silhouette_samples, silhouette_score
from sklearn.metrics.cluster import adjusted_rand_score
from sklearn.model_selection import RandomizedSearchCV
from sklearn.model_selection import RepeatedStratifiedKFold, KFold, cross_val_predict, train_test_split
from sklearn.model_selection import StratifiedKFold, GridSearchCV, cross_val_score, cross_validate

from sklearn.naive_bayes import GaussianNB
from sklearn.neighbors import KNeighborsRegressor, KNeighborsClassifier
from sklearn.pipeline import make_pipeline, Pipeline
from sklearn.preprocessing import MinMaxScaler, scale, StandardScaler, RobustScaler
from sklearn.preprocessing import LabelEncoder, OneHotEncoder, PolynomialFeatures, PowerTransformer
from sklearn.svm import SVR, SVC
from sklearn.tree import DecisionTreeClassifier, plot_tree

from catboost import CatBoostClassifier
from lightgbm import LGBMClassifier
from xgboost import XGBRegressor, XGBClassifier, plot_importance

# Ignore Warnings
import warnings
warnings.filterwarnings("ignore")
warnings.warn("this will not show")

# Figure&Display options
plt.rcParams["figure.figsize"] = (10,6)
pd.set_option('max_colwidth',200)
pd.set_option('display.max_rows', 1000)
pd.set_option('display.max_columns', 200)
pd.set_option('display.float_format', lambda x: '%.3f' % x)

C:\Users\aryaa\AppData\Local\Temp\ipykernel_18572\3784034954.py:3: DeprecationWarning: `import pandas_profiling` is going to be deprecated by April 1st. Please use `import ydata_profiling` instead.
import pandas_profiling
```

## 2.1 User Defined Functions

We have defined some useful user defined functions.

In [2]: ## Some Useful Functions

```
#####

def missing_values(df):
    missing_number = df.isnull().sum().sort_values(ascending = False)
    missing_percent = (df.isnull().sum() / df.isnull().count()).sort_values(ascending = False)
    missing_values = pd.concat([missing_number, missing_percent], axis = 1, keys = ['Missing_Number', 'Missing_Percent'])
    return missing_values[missing_values['Missing_Number'] > 0]

#####

def first_looking(df):
    print(colored("Shape:", attrs=['bold']), df.shape, '\n',
          colored(''*100, 'red', attrs = ['bold']),
          colored("\nInfo:\n", attrs = ['bold']), sep = '')
    print(df.info(), '\n',
          colored(''*100, 'red', attrs = ['bold']), sep = '')
    print(colored("Number of Uniques:\n", attrs = ['bold']), df.nunique(), '\n',
          colored(''*100, 'red', attrs = ['bold']), sep = '')
    print(colored("Missing Values:\n", attrs=['bold']), missing_values(df), '\n',
          colored(''*100, 'red', attrs = ['bold']), sep = '')
    print(colored("All Columns:", attrs = ['bold']), list(df.columns), '\n',
          colored(''*100, 'red', attrs = ['bold']), sep = '')

    df.columns = df.columns.str.lower().str.replace('&', '_').str.replace(' ', '_')
    print(colored("Columns after rename:", attrs = ['bold']), list(df.columns), '\n',
          colored(''*100, 'red', attrs = ['bold']), sep = '')
    print(colored("Columns after rename:", attrs = ['bold']), list(df.columns), '\n',
          colored(''*100, 'red', attrs = ['bold']), sep = '')
    print(colored("Descriptive Statistics \n", attrs = ['bold']), df.describe().round(2), '\n',
          colored(''*100, 'red', attrs = ['bold']), sep = '') # Gives a statistical breakdown of the data.
    print(colored("Descriptive Statistics (Categorical Columns) \n", attrs = ['bold']), df.describe(include = object).T, '\n',
          colored(''*100, 'red', attrs = ['bold']), sep = '') # Gives a statistical breakdown of the data.

def multicollinearity_control(df):
    feature = []
    collinear = []
    for col in df.corr().columns:
        for i in df.corr().index:
            if (abs(df.corr()[col][i]) > .9 and abs(df.corr()[col][i]) < 1):
                feature.append(col)
                collinear.append(i)
                print(colored(f"Multicollinearity alert in between:{col} - {i}",
                              "red", attrs = ['bold']), df.shape, '\n',
                      colored(''*100, 'red', attrs = ['bold']), sep = '')

def duplicate_values(df):
    print(colored("Duplicate check...", attrs = ['bold']), sep = '')
    print("There are", df.duplicated(subset = None, keep = 'first').sum(), "duplicated observations in the dataset.")
    duplicate_values = df.duplicated(subset = None, keep = 'first').sum()
    if duplicate_values > 0:
        df.drop_duplicates(keep = 'first', inplace = True)
        print(duplicate_values, colored(" Duplicates were dropped!"), '\n',
              colored(''*100, 'red', attrs = ['bold']), sep = '')
    # else:
    #     print(colored("There are no duplicates"), '\n',
    #           colored(''*100, 'red', attrs = ['bold']), sep = '')

# def drop_columns(df, drop_columns):
#     if drop_columns != []:
#         df.drop(drop_columns, axis = 1, inplace = True)
#         print(drop_columns, 'were dropped')
#     else:
#         print(colored('We will now check the missing values and if necessary, the related columns will be dropped!', attrs = ['bold']), '\n',
#               colored(''*100, 'red', attrs = ['bold']), sep = '')

def drop_null(df, limit):
    print('Shape:', df.shape)
    for i in df.isnull().sum().index:
        if (df.isnull().sum()[i] / df.shape[0]*100) > limit:
            print(df.isnull().sum()[i], 'percent of', i, 'null and were dropped')
            df.drop(i, axis = 1, inplace = True)
            print('new shape:', df.shape)
    print('New shape after missing value control:', df.shape)

#####

# To view summary information about the columns

def first_look(col):
    print("column name      : ", col)
    print("-----")
    print("Per_of_Nulls      : ", "%", round(df[col].isnull().sum() / df.shape[0]*100, 2))
    print("Num_of_Nulls       : ", df[col].isnull().sum())
    print("Num_of_Uniques      : ", df[col].nunique())
    print("Duplicates          : ", df.duplicated(subset = None, keep = 'first').sum())
    print(df[col].value_counts(dropna = False))

#####

def fill_most(df, group_col, col_name):
    '''Fills the missing values with the most existing value (mode) in the relevant column according to single-stage grouping'''
    for group in list(df[group_col].unique()):
        cond = df[group_col] == group
        mode = list(df[cond][col_name].mode())
        if mode != []:
            df.loc[cond, col_name] = df.loc[cond, col_name].fillna(df[cond][col_name].mode()[0])
        else:
            df.loc[cond, col_name] = df.loc[cond, col_name].fillna(df[col_name].mode()[0])
    print("Number of NaN : ", df[col_name].isnull().sum())
    print("-----")
    print(df[col_name].value_counts(dropna = False))

#####
# bar grafiğindeki değerlerin gösterilmesi
# show values in bar graphic
def show_values_on_bars(axes):
    def _show_on_single_plot(ax):
        for p in ax.patches:
            _x = p.get_x() + p.get_width() / 2
            _y = p.get_y() + p.get_height()
            value = '{:.2f}'.format(p.get_height())
            ax.text(_x, _y, value, ha="center")
    if isinstance(axes, np.ndarray):
        for idx, ax in np.ndenumerate(axes):
            _show_on_single_plot(ax)
    else:
        _show_on_single_plot(axes)
```

### 3 - ANALYSIS

### 3.1 Loading & Reading the Data

Let's first load the required HR dataset using pandas's "read\_csv" function.

```
In [3]: df0 = pd.read_csv('HR_Dataset.csv')
df = df0.copy()
df.head(3)
```

Out[3]:

	satisfaction_level	last_evaluation	number_project	average_monthly_hours	time_spend_company	Work_accident	left	promotion_last_5years	Departments	salary
0	0.380	0.530	2	157	3	0	1	0	sales	low
1	0.800	0.860	5	262	6	0	1	0	sales	medium
2	0.110	0.880	7	272	4	0	1	0	sales	medium

```
In [4]: df.tail(3)
```

Out[4]:

	satisfaction_level	last_evaluation	number_project	average_monthly_hours	time_spend_company	Work_accident	left	promotion_last_5years	Departments	salary
14996	0.370	0.530	2	143	3	0	1	0	support	low
14997	0.110	0.960	6	280	4	0	1	0	support	low
14998	0.370	0.520	2	158	3	0	1	0	support	low

```
In [5]: df.sample(3)
```

Out[5]:

	satisfaction_level	last_evaluation	number_project	average_monthly_hours	time_spend_company	Work_accident	left	promotion_last_5years	Departments	salary
13182	0.800	0.800	4	263	4	0	0	0	support	medium
8435	0.570	0.370	3	108	4	0	0	0	technical	low
5228	0.270	0.450	3	239	4	0	0	0	technical	low

## 4 - DATA CLEANING & EXPLORATORY DATA ANALYSIS (EDA)

Exploratory Data Analysis is an initial process of analysis, in which you can summarize characteristics of data such as pattern, trends, outliers, and hypothesis testing using descriptive statistics and visualization.

### 4.1 - A General Look at the Data

```
In [6]: first_looking(df)
duplicate_values(df)
print(colored("Shape:", attrs = ['bold']), df.shape, '\n', colored('***100, 'red', attrs = ['bold']))
```

```
Shape:(14999, 10)
*****
Info:

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14999 entries, 0 to 14998
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   satisfaction_level      14999 non-null  float64
1   last_evaluation        14999 non-null  float64
2   number_project         14999 non-null  int64
3   average_montly_hours   14999 non-null  int64
4   time_spend_company     14999 non-null  int64
5   Work_accident          14999 non-null  int64
6   left                   14999 non-null  int64
7   promotion_last_5years  14999 non-null  int64
8   Departments            14999 non-null  object
9   salary                 14999 non-null  object
dtypes: float64(2), int64(6), object(2)
memory usage: 1.1+ MB
None
*****
Number of Uniques:
satisfaction_level      92
last_evaluation         65
number_project          6
average_montly_hours    215
time_spend_company      8
Work_accident           2
left                    2
promotion_last_5years   2
Departments            10
salary                  3
dtype: int64
*****
Missing Values:
Empty DataFrame
Columns: [Missing_Number, Missing_Percent]
Index: []
*****
All Columns:['satisfaction_level', 'last_evaluation', 'number_project', 'average_montly_hours', 'time_spend_company', 'Work_accident', 'left', 'pro
motion_last_5years', 'Departments ', 'salary']
*****
Columns after rename:['satisfaction_level', 'last_evaluation', 'number_project', 'average_montly_hours', 'time_spend_company', 'work_accident', 'le
ft', 'promotion_last_5years', 'departments_', 'salary']
*****
Columns after rename:['satisfaction_level', 'last_evaluation', 'number_project', 'average_montly_hours', 'time_spend_company', 'work_accident', 'le
ft', 'promotion_last_5years', 'departments_', 'salary']
*****
Descriptive Statistics

      satisfaction_level  last_evaluation  number_project  \
count      14999.000      14999.000      14999.000
mean         0.610         0.720         3.800
std          0.250         0.170         1.230
min          0.090         0.360         2.000
25%          0.440         0.560         3.000
50%          0.640         0.720         4.000
75%          0.820         0.870         5.000
max          1.000         1.000         7.000

      average_montly_hours  time_spend_company  work_accident  left  \
count      14999.000      14999.000      14999.000  14999.000
mean       201.050         3.500         0.140     0.240
std        49.940         1.460         0.350     0.430
min        96.000         2.000         0.000     0.000
25%       156.000         3.000         0.000     0.000
50%       200.000         3.000         0.000     0.000
75%       245.000         4.000         0.000     0.000
max       310.000         10.000         1.000     1.000

      promotion_last_5years
count      14999.000
mean         0.020
std         0.140
min         0.000
25%         0.000
50%         0.000
75%         0.000
max         1.000
*****
Descriptive Statistics (Categorical Columns)
      count  unique  top  freq
departments_  14999    10  sales  4140
salary       14999     3   low  7316
*****
Duplicate check...
There are 3008 duplicated observations in the dataset.
3008 Duplicates were dropped!
*****
Shape: (11991, 10)
*****
```

**\*According to the basic examinations on the dataset:\***

- We have a classification problem.
- We are going to make classification on the target variable "left".
- And we will build a model to get the best classification on the "left" column.
- Because of that we are going to look at the balance of "left" column.
- The dataset has 10 columns and 11991 observations after dropping of duplicated observations.
- 8 columns contain numerical values and 2 columns contain categorical values.
- There seems to be no missing value.

```
In [7]: df.columns

Out[7]: Index(['satisfaction_level', 'last_evaluation', 'number_project',
              'average_montly_hours', 'time_spend_company', 'work_accident', 'left',
              'promotion_last_5years', 'departments_', 'salary'],
              dtype='object')

In [8]: df.rename({'departments_': 'department'}, axis=1, inplace=True)
df.head(1)

Out[8]:   satisfaction_level  last_evaluation  number_project  average_montly_hours  time_spend_company  work_accident  left  promotion_last_5years  department  salary
0          0.380          0.530           2          157           3           0    1           0      sales      low

In [9]: df = df[['satisfaction_level', 'last_evaluation', 'number_project',
              'average_montly_hours', 'time_spend_company', 'work_accident',
              'promotion_last_5years', 'department', 'salary', 'left']]
df.head(1)

Out[9]:   satisfaction_level  last_evaluation  number_project  average_montly_hours  time_spend_company  work_accident  promotion_last_5years  department  salary  left
0          0.380          0.530           2          157           3           0           0      sales      low    1
```

- I want to move the 'left' column, which is my target column, from where it is to the end. In this way, I will work more comfortably psychologically :))

## 4.2 - Examination of Features and Data Insights

In the given dataset, we have two types of employee one who stayed and another who left the company. So, we can divide data into two groups and compare their characteristics. Here, we can find the average of both the groups using groupby() and mean() function.


### 'Left' Column-Target Column

```
In [10]: cprint("Have a First Look to 'left' Column",'green')
first_look('left')

Have a First Look to 'left' Column
column name      : left
-----
Per_of_Nulls     : % 0.0
Num_of_Nulls     : 0
Num_of_Uniques   : 2
Duplicates       : 0
0      10000
1      1991
Name: left, dtype: int64

In [11]: import plotly
import plotly.express as px
fig = px.pie(df, values = df['left'].value_counts(),
            names = (df['left'].value_counts()).index,
            title = "'left' Column Distribution")
fig.show()
```

"left" Column Distribution



```
In [12]: y = df['left']
print(f'Percentage of left-1: % {round(y.value_counts(normalize=True)[1]*100,2)} --> \
({y.value_counts()[1]} observations for left-1)\nPercentage of left-0: % {round(y.value_counts(normalize=True)[0]*100,2)} --> ({y.value_counts()[0]}
observations for left-0)')

Percentage of left-1: % 16.6 --> (1991 observations for left-1)
Percentage of left-0: % 83.4 --> (10000 observations for left-0)

• 'left' column has binary type values.
• We have an imbalanced data.
• Almost 17% of the employees didn't continue with the company and left.
• 1991 employees left.
• Almost 83% of the employees continue with the company and didn't leave.
• 10000 employees didn't leave.
```

```
In [13]: df.groupby('left').mean()
```

	satisfaction_level	last_evaluation	number_project	average_monthly_hours	time_spend_company	work_accident	promotion_last_5years
left							
0	0.667	0.716	3.787	198.943	3.262	0.174	0.019
1	0.440	0.722	3.883	208.162	3.881	0.053	0.004

```
In [14]: cprint('Dataset describe results according to the "left==1" condition','green', 'on_black')
df[df['left'] == 1].describe().T.style.background_gradient(subset = ['mean','min','50%', 'max'], cmap = 'RdPu')
```

Dataset describe results according to the "left==1" condition

	count	mean	std	min	25%	50%	75%	max
satisfaction_level	1991.000000	0.440271	0.265207	0.090000	0.110000	0.410000	0.730000	0.920000
last_evaluation	1991.000000	0.721783	0.197436	0.450000	0.520000	0.790000	0.910000	1.000000
number_project	1991.000000	3.883476	1.817139	2.000000	2.000000	4.000000	6.000000	7.000000
average_monthly_hours	1991.000000	208.162230	61.295145	126.000000	146.000000	226.000000	262.500000	310.000000
time_spend_company	1991.000000	3.881467	0.974041	2.000000	3.000000	4.000000	5.000000	6.000000
work_accident	1991.000000	0.052737	0.223565	0.000000	0.000000	0.000000	0.000000	1.000000
promotion_last_5years	1991.000000	0.004018	0.063277	0.000000	0.000000	0.000000	0.000000	1.000000
left	1991.000000	1.000000	0.000000	1.000000	1.000000	1.000000	1.000000	1.000000

```
In [15]: cprint('Dataset describe results according to the "left==0" condition','green', 'on_black')
df[df['left'] == 0].describe().T.style.background_gradient(subset = ['mean','min','50%', 'max'], cmap = 'RdPu')
```

Dataset describe results according to the "left==0" condition

	count	mean	std	min	25%	50%	75%	max
satisfaction_level	10000.000000	0.667365	0.217082	0.120000	0.540000	0.690000	0.840000	1.000000
last_evaluation	10000.000000	0.715667	0.161919	0.360000	0.580000	0.710000	0.850000	1.000000
number_project	10000.000000	3.786800	0.981755	2.000000	3.000000	4.000000	4.000000	6.000000
average_monthly_hours	10000.000000	198.942700	45.665507	96.000000	162.000000	198.000000	238.000000	287.000000
time_spend_company	10000.000000	3.262000	1.367239	2.000000	2.000000	3.000000	4.000000	10.000000
work_accident	10000.000000	0.174500	0.379558	0.000000	0.000000	0.000000	0.000000	1.000000
promotion_last_5years	10000.000000	0.019500	0.138281	0.000000	0.000000	0.000000	0.000000	1.000000
left	10000.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000

### 'satisfaction\_level' Column


```
In [16]: print("Have a First Look to 'left' Column")
first_look('satisfaction_level')

Have a First Look to 'left' Column
column name      :  satisfaction_level
-----
Per_of_Nulls    :   % 0.0
Num_of_Nulls    :    0
Num_of_Uniques  :   92
Duplicates      :    0
0.740      214
0.100      203
0.730      201
0.500      200
0.720      199
0.840      199
0.830      196
0.770      194
0.780      194
0.660      192
0.890      191
0.110      188
0.750      188
0.760      188
0.910      187
0.800      186
0.590      185
0.630      184
0.820      180
0.570      179
0.960      179
0.690      178
0.850      177
0.490      177
0.790      177
0.810      176
0.610      176
0.870      176
0.700      176
0.900      172
0.520      172
0.650      171
0.860      170
0.600      167
0.560      166
0.540      166
0.980      165
0.640      164
0.920      163
0.580      162
0.620      162
0.880      160
0.510      160
0.550      159
0.670      158
0.990      157
0.950      153
0.710      152
0.970      150
0.530      150
0.930      148
0.680      145
0.940      143
0.430      141
0.440      138
0.370      130
0.400      126
0.450      121
0.480      120
0.380      116
0.390      111
0.090      110
0.420      108
0.410      104
1.000      100
0.360       98
0.150       69
0.240       68
0.460       63
0.210       63
0.140       62
0.170       62
0.160       59
0.190       58
0.180       56
0.200       56
0.220       53
0.230       51
0.130       45
0.340       44
0.310       42
0.320       42
0.300       38
0.350       36
0.470       34
0.290       32
0.330       30
0.250       29
0.260       28
0.120       26
0.280       24
0.270       23
Name: satisfaction_level, dtype: int64

In [17]: pd.crosstab(df['satisfaction_level'], df['left']).iplot(kind='bar', title = 'satisfaction_level and left')
```



satisfaction\_level and left



[Export to plot.ly »](#)

- Although it comes to mind that there should be a linear relationship between 'satisfaction\_level' and 'left', it does not look like this on the graph.
- Those with a 'satisfaction\_level' value of around 0.1 are very likely to 'left'.
- There is a significant increase in the number of those whose 'satisfaction\_level' value is between 3.5 and 4.5 and 'left'. In fact, the number of left ones exceeds the notleft ones.
- When the 'satisfaction\_level' value is between 7 and 9, there is an increase in the number of those left.
- Normally we expect low satisfaction level for the employees who has left, so the part near to 0 on the x-axis is make sense.
- Besides a group of employee who are not very decisive about their satisfaction level have also been left the company. This group may need extra motivation for employee loyalty. Because they are not so clear in their assesments about their future in the company.
- Also a group of employee whose satisfaction level is above the avarage have been left the company. This does not make sense so this must be investigated deeply.

There may be some other issues:

- a. The method of gathering this information may be wrong. So the assessment of satisfaction level and the resignings may not be directly proportional.
- b. The assessment may not be up to date. By the time the satisfaction level may be decreased so at the real time the satisfaction level of all resigning employees may be close to 0.
- c. Some of the employees may have hidden their true feelings.


**'last\_evaluation' Column**

```
In [18]: cprint("Have a First Look to 'last_evaluation' Column",'green', 'on_black')
         first_look('last_evaluation')
```

```
Have a First Look to 'last_evaluation' Column
column name      : last_evaluation
-----
Per_of_Nulls     : % 0.0
Num_of_Nulls     : 0
Num_of_Uniques   : 65
Duplicates       : 0
0.550            281
0.500            269
0.510            264
0.570            258
0.540            252
0.560            248
0.530            247
0.850            244
0.900            244
0.490            242
0.870            241
0.890            235
0.520            233
0.740            227
0.910            224
0.590            221
0.860            219
0.840            218
0.970            218
0.770            212
0.480            212
0.830            211
0.950            211
0.810            210
0.750            209
0.930            209
0.960            208
0.980            208
0.920            207
0.670            205
0.630            204
0.580            203
0.800            201
1.000            200
0.610            200
0.940            199
0.640            198
0.620            197
0.990            197
0.790            195
0.820            195
0.680            193
0.760            192
0.660            191
0.730            189
0.880            187
0.600            186
0.700            184
0.720            182
0.780            181
0.650            176
0.690            176
0.710            171
0.460            131
0.470            121
0.450            78
0.410            52
0.400            49
0.370            46
0.390            45
0.420            45
0.430            44
0.380            42
0.440            35
0.360            19
Name: last_evaluation, dtype: int64
```

```
In [19]: pd.crosstab(df['last_evaluation'], df['left']).plot(kind='bar', title = 'last_evaluation and left')
```

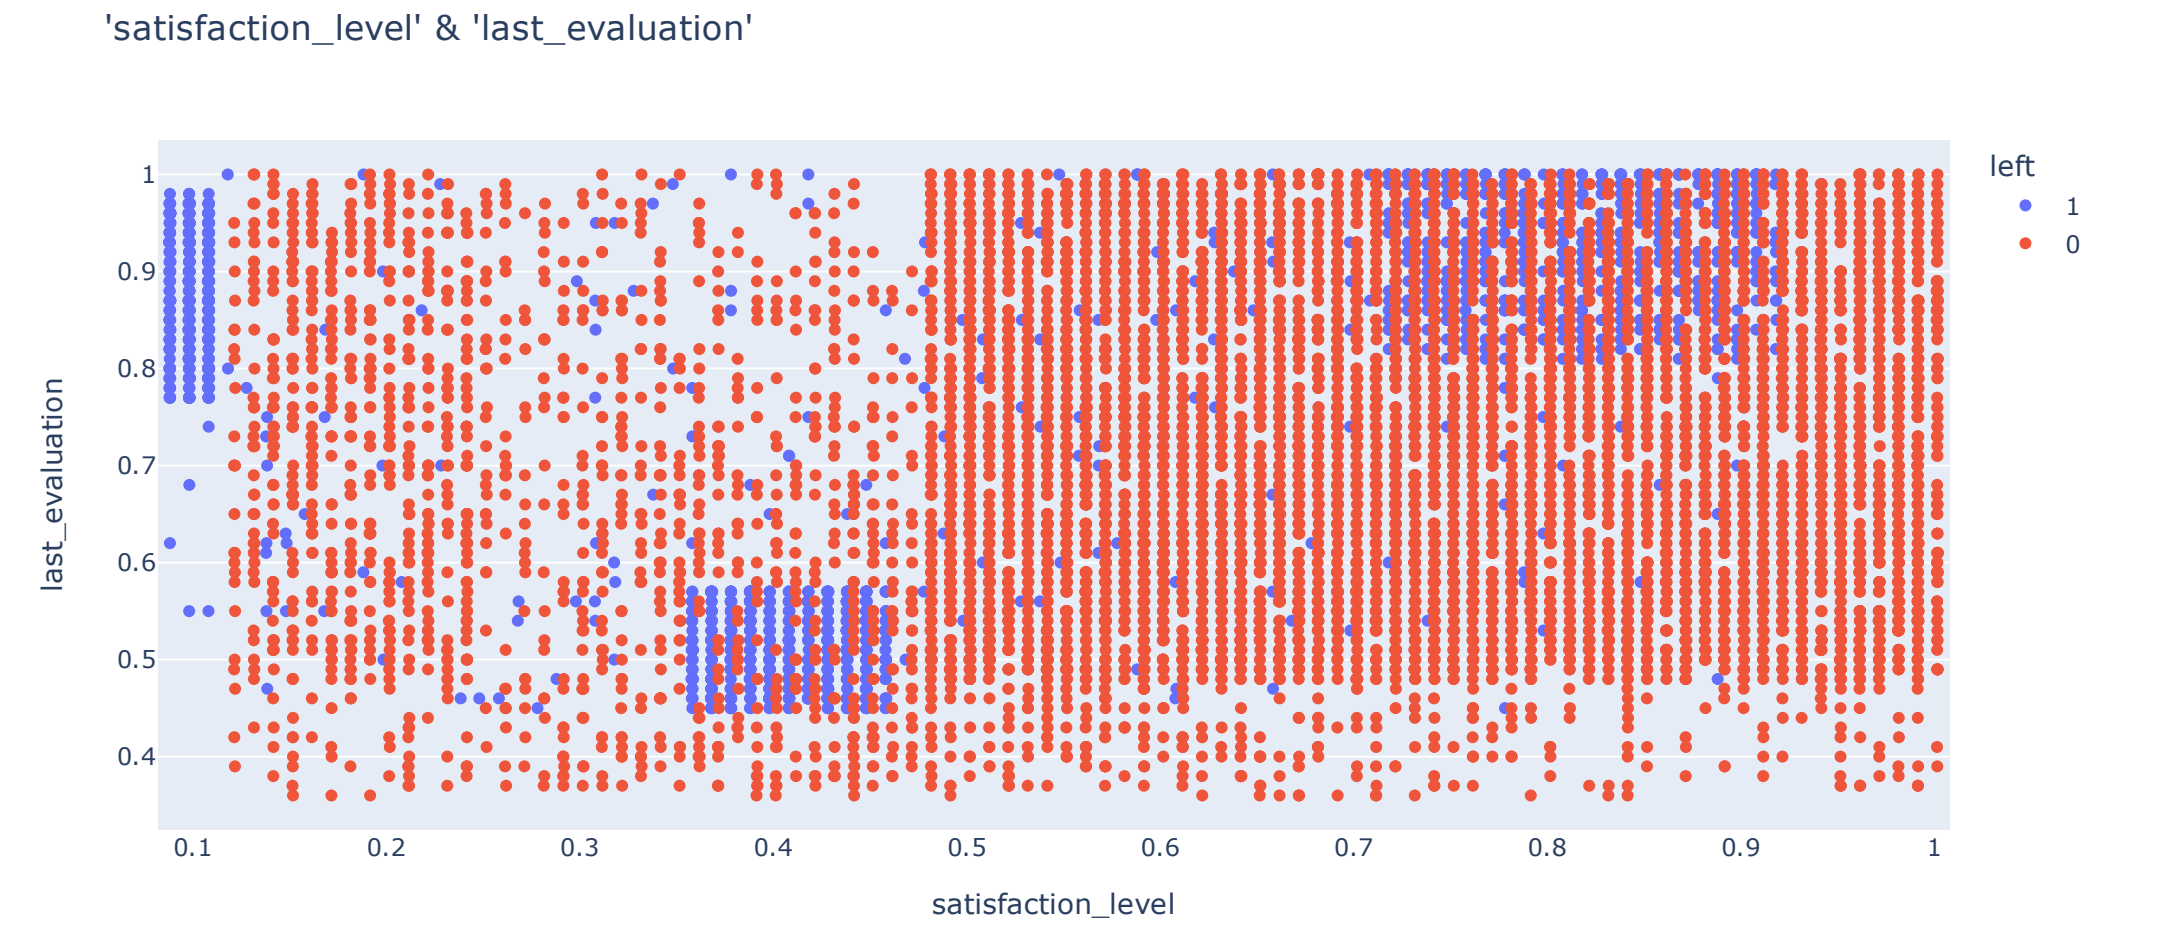
last\_evaluation and left



[Export to plot.ly »](#)

- Most of the employees have been assessed above 0.4.
- There is a local increase between 0.45-0.6 and 0.8-1 in 'last\_evaluation' values, as in 'satisfaction\_level' values. There is an increase in the number of people who quit their jobs in these intervals
- Intensive work may cause the resign of high evaluated employees (second group). Because employer will be happy with performance of these staff, however it will be a burden for employee.

```
In [20]: fig = px.strip(df, x = 'satisfaction_level', y = 'last_evaluation', color = 'left',
                    title = "'satisfaction_level' & 'last_evaluation'")
fig.show()
```



It becomes meaningful when the satisfaction level of employees and the evaluation of the employer shown together.

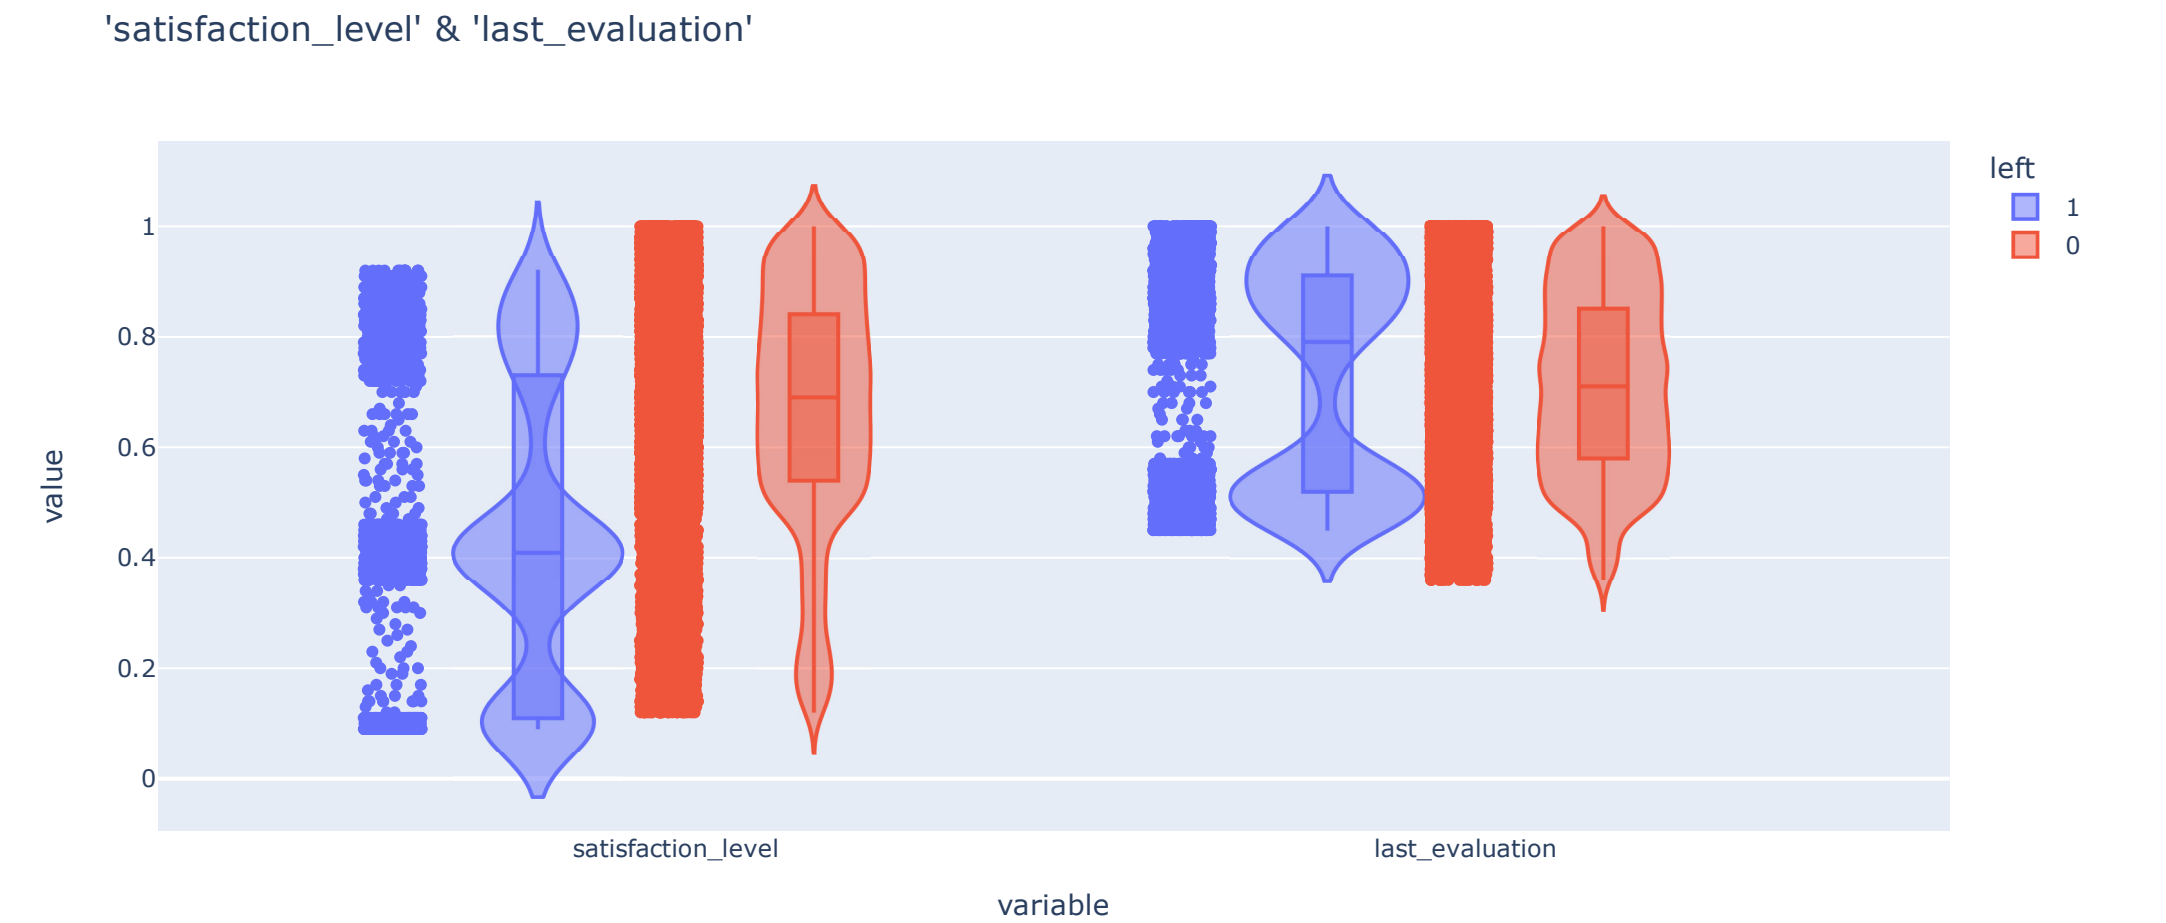
As seen in the graph; the resigning employees are grouping in three diferent clusters.

1. First group has a satisfaction level of 0.4 and last evaluation of 0.5. This group has not a clear idea about the company and the employer does not have a clear assessment about them. Other features affecting this group must to be investigated. What are the main questions of this group? Why they are confusing? What are the pros and cons of the company for these group? and so on...
2. The second group has a low satisfaction even if the employer evaluated them with high degrees. Then what can be the main problem of this group?

Intensive work with a low salary may affect this group. Or intensive work without promotion may cause to leave. On the next steps workload and motivation factors of this group have to be investigated.

3. The third group has a high satisfaction level and evaluation point as well. The density of this group is fewer than the others. The issues that triger the leave of this group need to be investigated.

```
In [21]: fig = px.violin(df[['satisfaction_level', 'last_evaluation', 'left']], color = 'left', box = True, points='all',
                    title = "'satisfaction_level' & 'last_evaluation'")
fig.show()
```




#### 'number\_project' Column

```
In [22]: cprint("Have a First Look to 'number_project' Column",'green', 'on_black')
first_look('number_project')
```

```
Have a First Look to 'number_project' Column
column name   : number_project
-----
Per_of_Nulls  : % 0.0
Num_of_Nulls  : 0
Num_of_Uniques : 6
Duplicates    : 0
4    3685
3    3520
5    2233
2    1582
6     826
7     145
Name: number_project, dtype: int64
```


```
In [23]: fig = px.pie(df, values = df['number_project'].value_counts(),
                    names = (df['number_project'].value_counts()).index,
                    title = "'number_project' Column Distribution")
fig.show()
```

"number\_project" Column Distribution



```
In [24]: pd.crosstab(df['number_project'], df['left']).plot(kind='bar', title = 'number_project and left')
```

number\_project and left



[Export to plot.ly »](#)


The number of leaving employees is higher among those who have only two projects during the period. This can be summed up as: "the employees with only two projects feel worthless or emptied". Because most of the employees work on three or four projects.

With the 6th project, the number of resignings is getting over the number of ongoing. There are no ongoing staff members who were assigned to 7 projects.

Working on more projects may cause intensive workload, regarding to this the satisfaction level may decrease with the insufficient motivators.

```
In [25]: fig = px.strip(df[df['number_project'] == 2], x = 'satisfaction_level', y = 'last_evaluation', color = 'left',
                  title = "'satisfaction_level' & 'last_evaluation' when 'number_project' == 2")
fig.show()
```

'satisfaction\_level' & 'last\_evaluation' when 'number\_project' == 2




If we look at the satisfaction level, evaluation score and the number of projects together;

The group of undecideds who were evaluated as 0.5 are the group who worked on only two projects. As a result, our hypothesis about this group is becoming more clear. As the employer does not assign enough projects to this group, he/she cannot evaluate their performance and they feel worthless. Therefore, they are unsure about their future in the company. This may lead them to leave.


```
In [26]: fig = px.strip(df[df['number_project'] > 4], x = 'satisfaction_level', y = 'last_evaluation', color = 'left',
                  title = "'satisfaction_level' & 'last_evaluation' when 'number_project' > 4")
fig.show()
```

'satisfaction\_level' & 'last\_evaluation' when 'number\_project' > 4



```
In [27]: fig = px.strip(df[df['number_project'] == 7], x = 'satisfaction_level', y = 'last_evaluation', color = 'left',
                    title = "'satisfaction_level' & 'last_evaluation' when 'number_project' == 7")
fig.show()
```

'satisfaction\_level' & 'last\_evaluation' when 'number\_project' == 7



The leaving employees who worked on more than four projects are the group two and three of the last\_evaluation section.

Especially most of the second group of last\_evaluation section are worked on seven projects and left the company. So again our hypothesis about this group is now more definite.

**'average\_montly\_hours' Column**

```
In [28]: cprint("Have a First Look to 'average_montly_hours' Column",'green', 'on_black')
first_look('average_montly_hours')
```


```
Have a First Look to 'average_montly_hours' Column
column name      : average_montly_hours
-----
Per_of_Nulls    : % 0.0
Num_of_Nulls    : 0
Num_of_Uniques  : 215
Duplicates      : 0
156      112
149      112
160      111
151      107
135      104
260      102
257      100
145      98
157      96
224      94
152      94
143      94
140      93
155      93
139      92
137      92
141      91
243      91
245      91
238      91
158      90
154      90
148      90
159      89
264      87
150      87
142      87
258      86
134      85
271      85
162      85
255      85
147      85
153      85
266      84
146      84
269      83
254      83
246      82
253      82
167      81
263      81
168      81
185      81
136      81
247      81
173      81
237      79
250      79
226      79
192      79
251      78
144      78
233      78
261      78
201      77
229      77
270      77
232      77
248      76
169      76
242      75
249      75
274      75
206      74
171      74
225      74
272      74
239      74
259      73
189      73
183      73
244      73
133      73
223      73
265      73
202      73
181      73
198      72
217      72
178      72
267      72
273      72
240      71
180      71
256      71
241      70
184      70
268      70
138      70
252      69
176      69
199      69
161      69
177      69
222      69
165      69
218      68
216      68
219      68
196      68
175      67
182      67
132      67
166      67
205      67
231      66
212      66
190      66
221      66
191      66
174      66
262      65
163      65
214      65
227      65
236      65
210      64
197      64
213      64
186      64
203      64
164      63
188      62
211      62
```

170	62
179	62
208	62
235	61
193	61
207	61
204	61
234	61
187	60
172	60
195	60
209	60
194	60
215	60
228	60
275	56
220	55
200	53
230	52
131	52
128	45
127	42
129	41
130	38
286	34
280	25
285	24
281	24
279	23
104	23
113	23
276	22
278	22
111	22
282	21
98	20
121	20
287	20
126	19
284	19
123	18
301	17
100	17
106	17
283	17
109	16
101	16
105	16
108	16
125	16
114	15
102	14
277	14
117	14
103	14
310	14
115	13
97	13
124	12
99	11
112	10
307	10
107	10
120	10
296	10
308	10
110	10
122	10
291	9
116	9
119	9
294	9
305	8
304	8
290	8
118	8
306	8
295	8
289	7
293	7
309	7
292	7
300	6
96	6
288	6
298	5
302	5
297	5
299	5
303	5

Name: average\_monthly\_hours, dtype: int64

In [29]: `pd.crosstab(df['average_monthly_hours'], df['left']).plot(kind='bar', title = 'average_monthly_hours and left')`

average\_monthly\_hours and left



[Export to plot.ly »](#)

- Looking at the 'average\_monthly\_hours' values, there is a local increase in turnover in the 125-160 month working hours range and 210-290 monthly working hours.
- Those who work more than 290 hours per month are more likely to quit their jobs than those who do not.
- So the next question is "The average monthly working hours are related to projects number or not?"



```
In [30]: plt.figure(figsize = (16,6))
sns.lineplot(data = df, x = 'average_monthly_hours', y = 'number_project', hue = 'left')
plt.title("'average_monthly_hours' & 'number_project'");
```

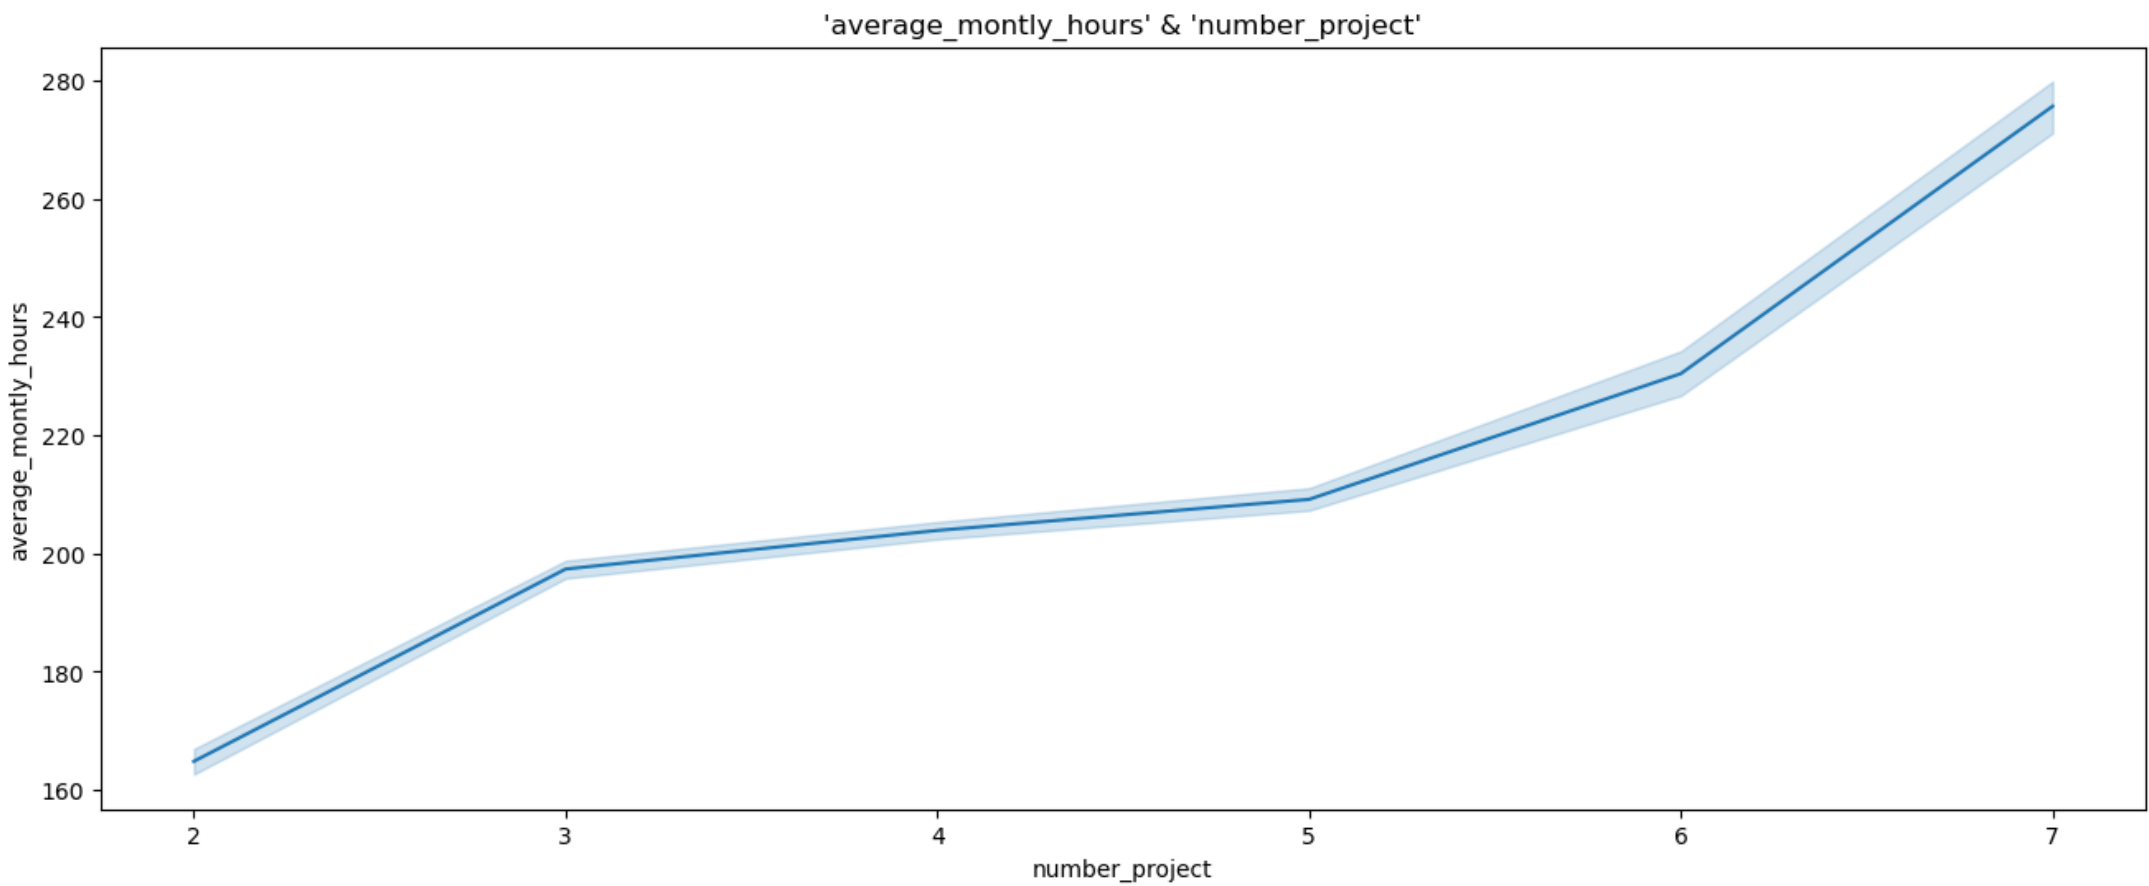


At the graph above it is seen that the group working on two projects is working nearly 130-160 hours monthly. It can be assessed that they have only two simple projects that they don't need to work hard, so their loyalty is weak.

Most of the employees are working 135-275 hours monthly. In this group usually the employees who get two or more than five projects leaving the company.

When there is an increase on the number of projects and the average monthly working hours, there is also an increase on the number of resignings.

```
In [31]: plt.figure(figsize = (16,6))
sns.lineplot(data = df, y = 'average_monthly_hours', x = 'number_project')
plt.title("'average_monthly_hours' & 'number_project'");
```



The increasing of the average monthly hours according to number of projects is seen on the graph.

The rate of increase is higher between two and three projects, and after five projects. So it is clearly define the number of resignings due to the working hours.

**There need to be an adjustment about the project numbers, working hours and workload. The projects must be assigned to more employees. Also, better incentives must be offered to staff who are working hard.**


#### 'time\_spend\_company' Column

```
In [32]: cprint("Have a First Look to 'time_spend_company' Column",'green', 'on_black')
first_look('number_project')
```

```
Have a First Look to 'time_spend_company' Column
column name      : number_project
-----
Per_of_Nulls    : % 0.0
Num_of_Nulls    : 0
Num_of_Uniques  : 6
Duplicates      : 0
4      3685
3      3520
5      2233
2      1582
6       826
7       145
Name: number_project, dtype: int64
```

```
In [33]: df['time_spend_company'].value_counts().plot(kind="bar", title = "'time_spend_company' Column Distribution')
```


"time\_spend\_company" Column Distribution






```
In [34]: fig = px.pie(df, values = df['time_spend_company'].value_counts(),
names = (df['time_spend_company'].value_counts()).index,
title = '"time_spend_company" Column Distribution')
fig.show()
```

"time\_spend\_company" Column Distribution



```
In [35]: pd.crosstab(df['time_spend_company'], df['left']).plot(kind='bar', title = 'time_spend_company and left')
```

time\_spend\_company and left

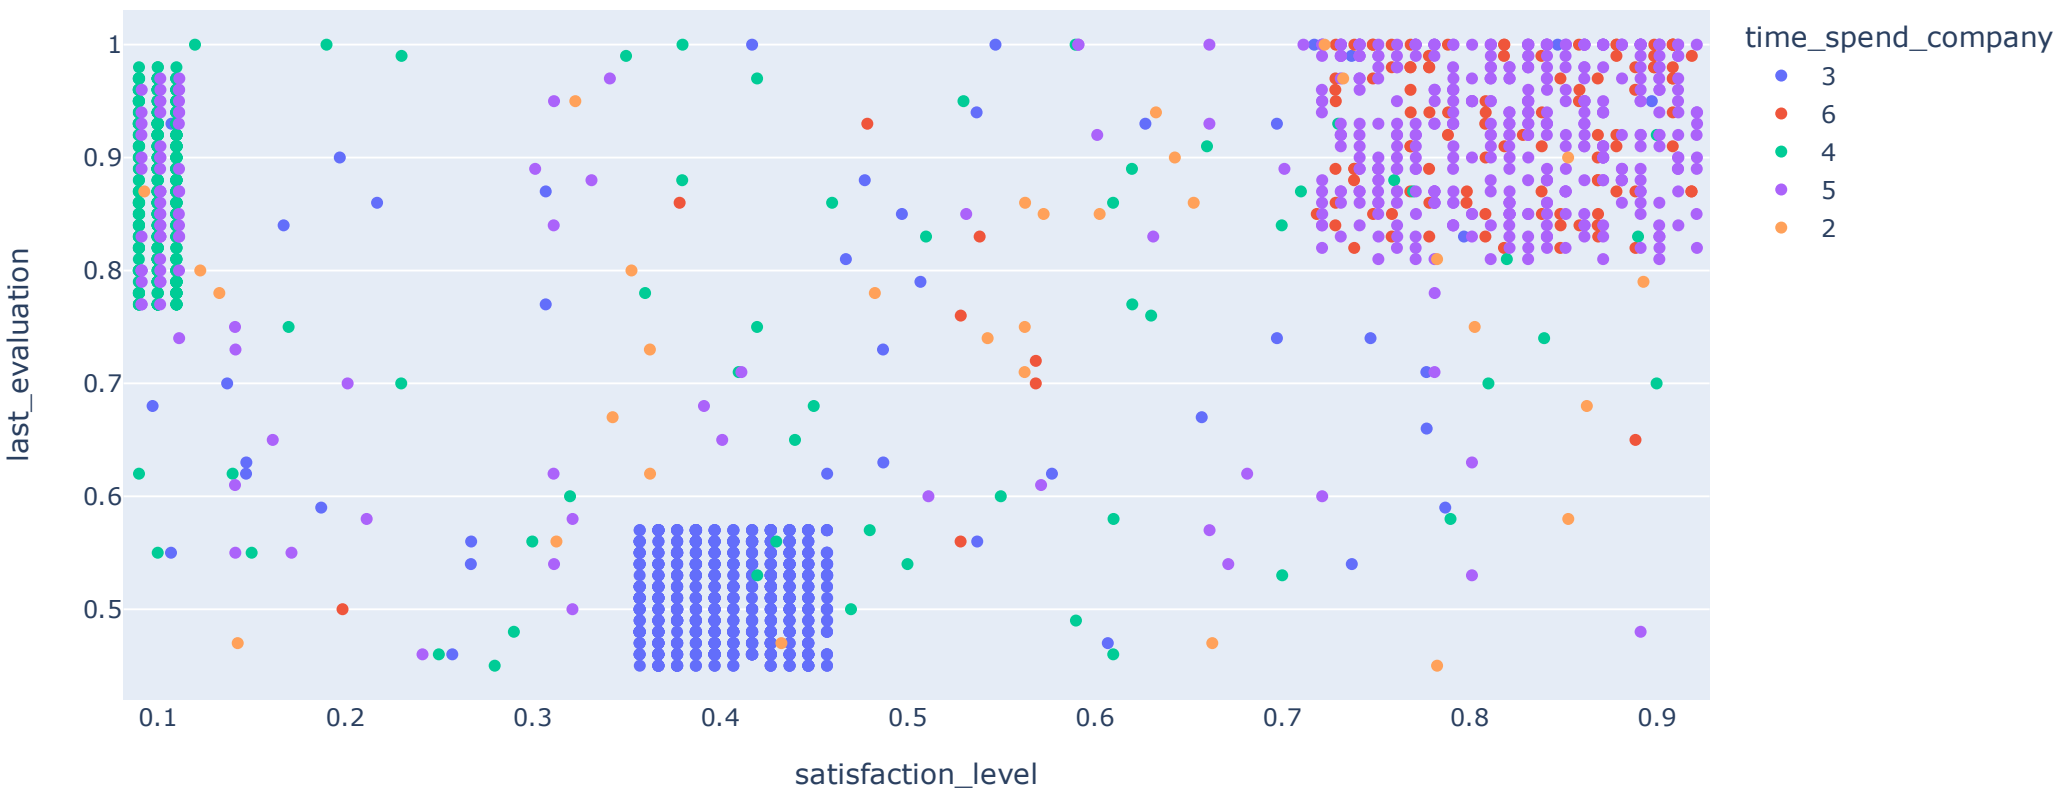


[Export to plot.ly »](#)

- Looking at the 'time\_spent\_company' values, there is an increase in turnover in the 3rd working year, but this increase gradually decreases until the 6th working year.

```
In [36]: fig = px.strip(df[df['left'] == 1], x = 'satisfaction_level', y = 'last_evaluation', color = 'time_spend_company',
title = '"satisfaction_level' & 'last_evaluation"')
fig.show()
```

'satisfaction\_level' & 'last\_evaluation'



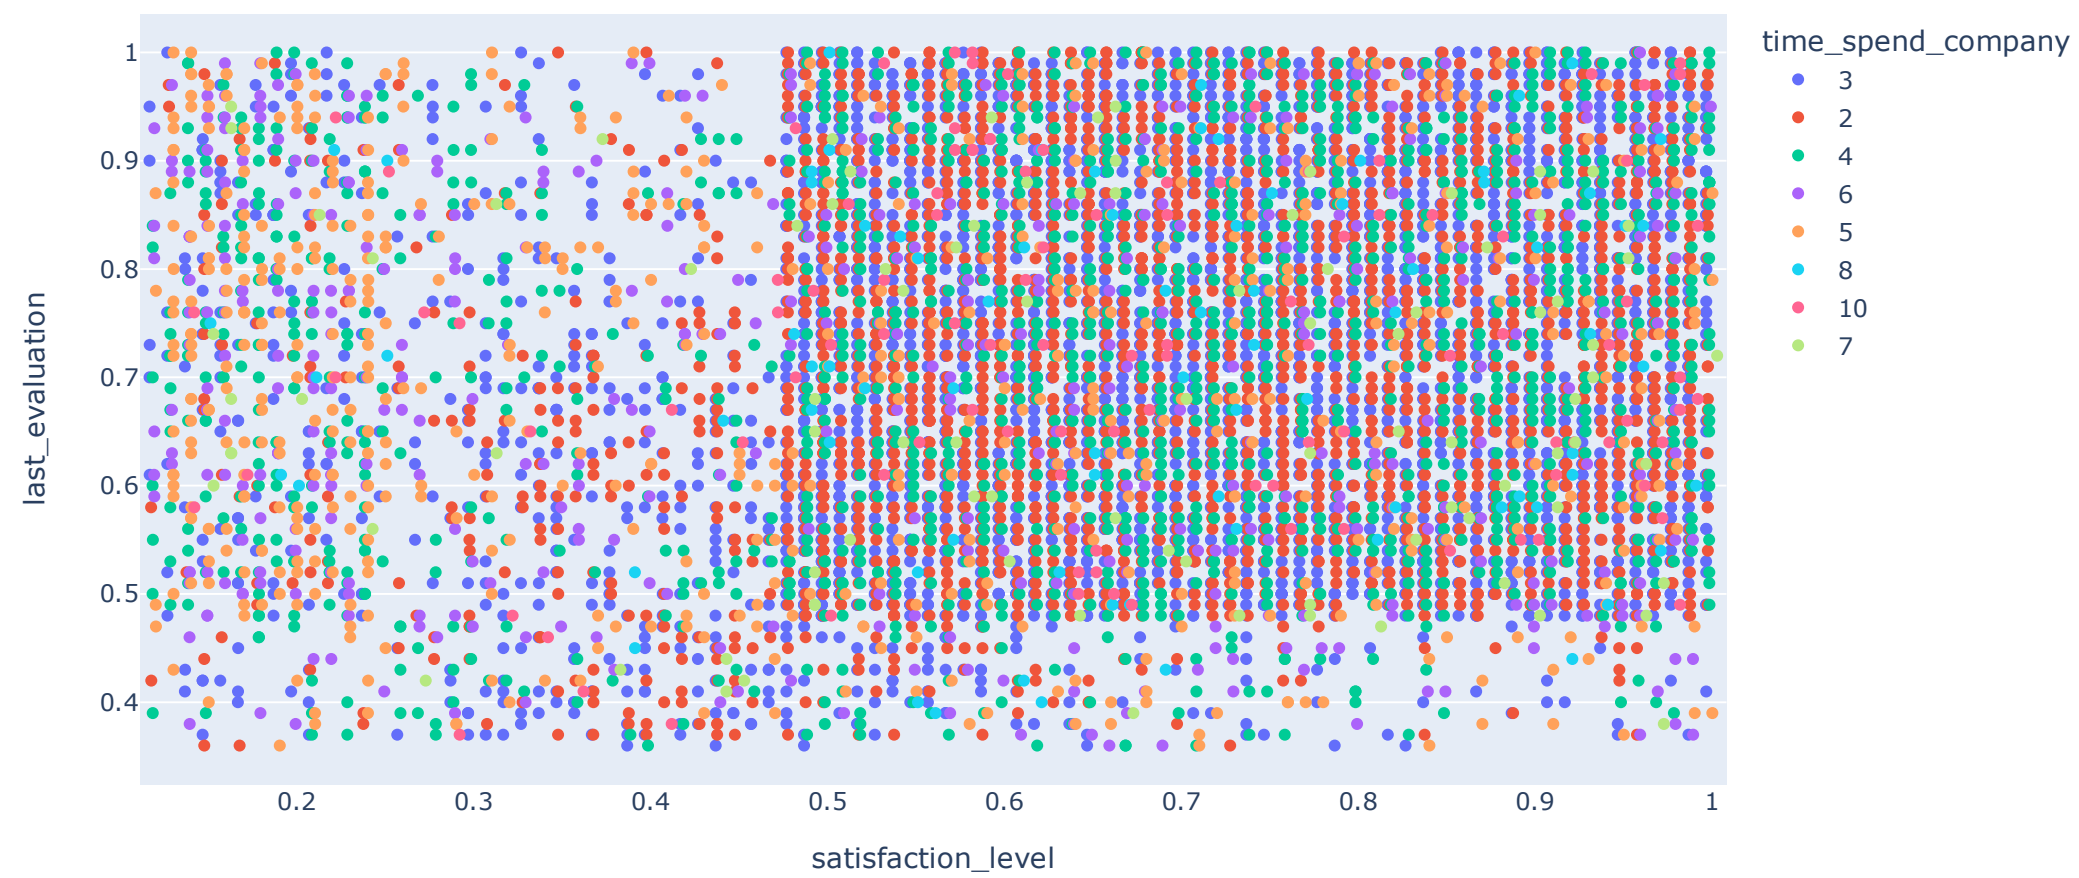
As can be seen on the graph, the employees are not able to make a clear assessment of the company during the first three years of their employment. This, coupled with the other factors, tends to lead to leaving the company after three years.

By the fourth year, their workload increases and their satisfaction decreases.

After the fifth year, they make an assessment, "they will leave or not".

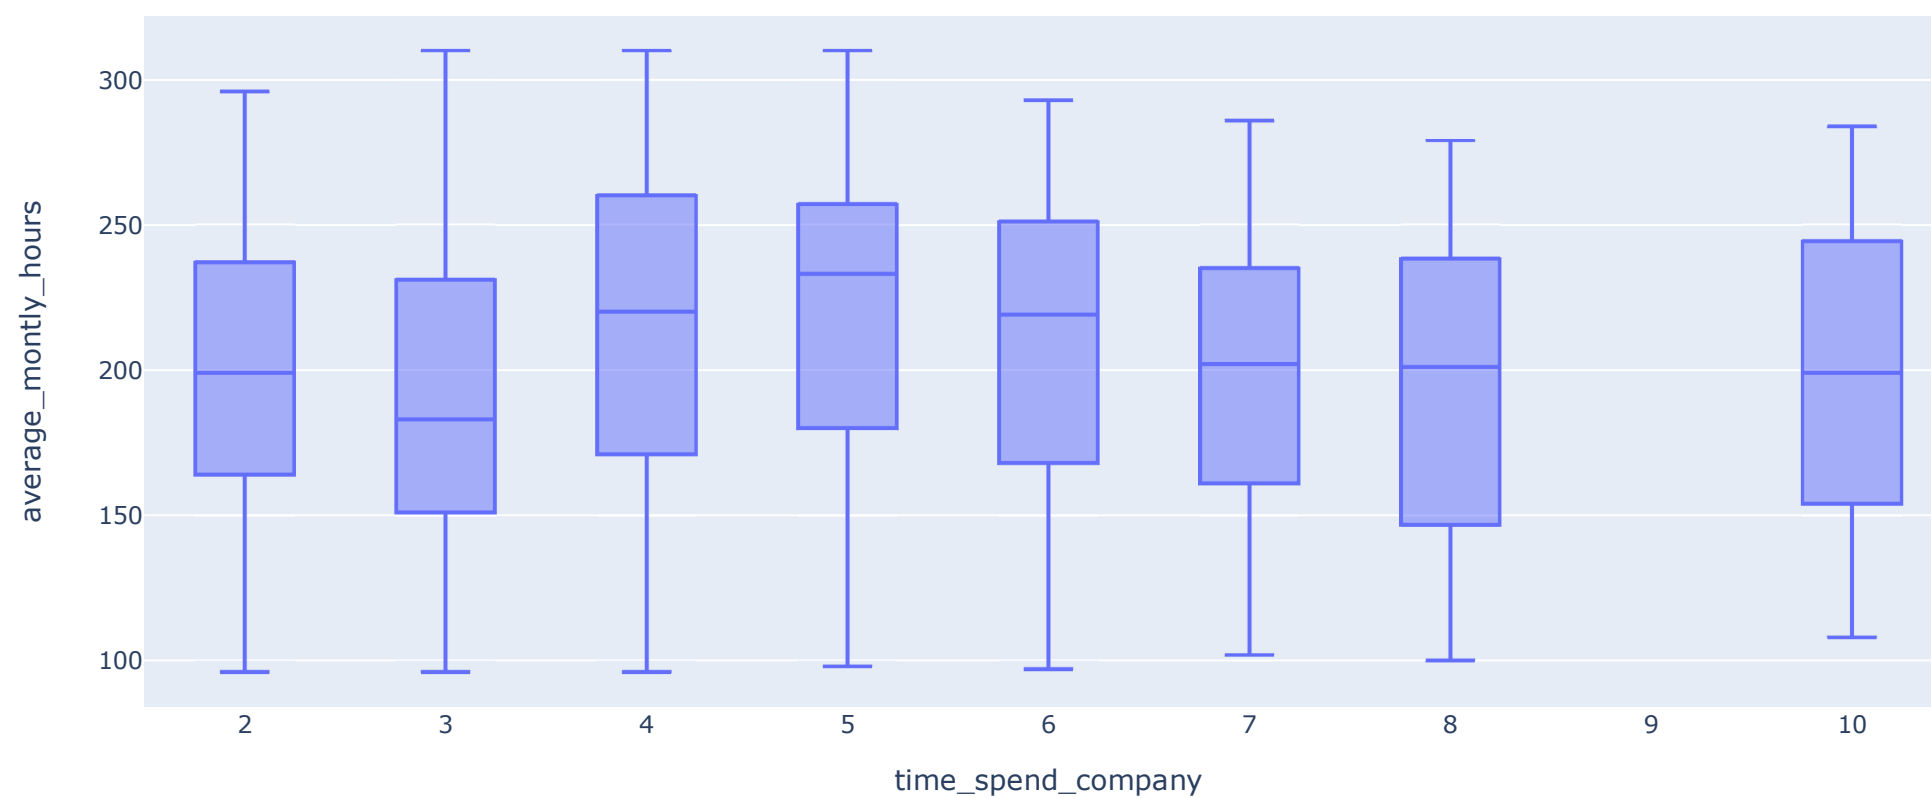
If they decide to continue in the company, they never consider leaving after the sixth year.

```
In [37]: fig = px.strip(df[df['left'] == 0], x = 'satisfaction_level', y = 'last_evaluation', color = 'time_spend_company')
fig.show()
```




```
In [38]: fig = px.box(df, x = 'time_spend_company', y = 'average_monthly_hours', title = "'time_spend_company' & 'average_monthly_hours'")
fig.show()
```

'time\_spend\_company' & 'average\_monthly\_hours'




```
In [39]: fig = px.box(df, x = 'time_spend_company', y = 'number_project', title = "'time_spend_company' & 'number_project'")
fig.show()
```

'time\_spend\_company' & 'number\_project'



```
In [40]: px.histogram(df, x = 'time_spend_company', color = 'number_project', title = "'time_spend_company' & 'number_project'")
```

'time\_spend\_company' & 'number\_project'



**\*Then how is the relation between workload and time spend in the company?\***

Third year staff has the most workload. After that year number of participated project is decreasing stepped. It makes sense. The experienced staff becoming team leader or manager position. That's why less of them can be assigned to projects.


**'work\_accident' Column**

```
In [41]: cprint("Have a First Look to 'work_accident' Column",'green', 'on_black')
first_look('work_accident')
```

```
Have a First Look to 'work_accident' Column
column name      : work_accident
-----
Per_of_Nulls     : % 0.0
Num_of_Nulls     : 0
Num_of_Uniques   : 2
Duplicates       : 0
0      10141
1      1850
Name: work_accident, dtype: int64
```


```
In [42]: fig = px.pie(df, values = df['work_accident'].value_counts(),
names = (df['work_accident'].value_counts()).index,
title = '"work_accident" Column Distribution')
fig.show()
```

"work\_accident" Column Distribution



```
In [43]: pd.crosstab(df['work_accident'], df['left']).iplot(kind='bar', title = 'work_accident and left')
```


work\_accident and left



[Export to plot.ly »](#)


```
In [44]: px.histogram(df, x = df['average_monthly_hours'], color='work_accident', title = 'work_accident and average_monthly_hours')
```

work\_accident and average\_monthly\_hours



```
In [45]: px.histogram(df, x = df['time_spend_company'], color='work_accident', title = 'work_accident and time_spend_company')
```

work\_accident and time\_spend\_company



- 'work\_accident' column has binary type values.
- Left ratios are similar between those who have had a work accident and those who have not.
- It does not appear to be a determining factor. In fact, it can be said that the left rate of those who have had a work accident is proportionally lower.


'promotion\_last\_5years' Column

```
In [46]: cprint("Have a First Look to 'promotion_last_5years' Column",'green', 'on_black')
first_look('promotion_last_5years')
```

```
Have a First Look to 'promotion_last_5years' Column
column name      : promotion_last_5years
-----
Per_of_Nulls     : % 0.0
Num_of_Nulls     : 0
Num_of_Uniques   : 2
Duplicates       : 0
0      11788
1       203
Name: promotion_last_5years, dtype: int64
```


```
In [47]: fig = px.pie(df, values = df['promotion_last_5years'].value_counts(),
names = (df['promotion_last_5years'].value_counts()).index,
title = '"promotion_last_5years" Column Distribution')
fig.show()
```

"promotion\_last\_5years" Column Distribution



```
In [48]: pd.crosstab(df['promotion_last_5years'], df['left']).iplot(kind='bar', title = 'promotion_last_5years and left')
```

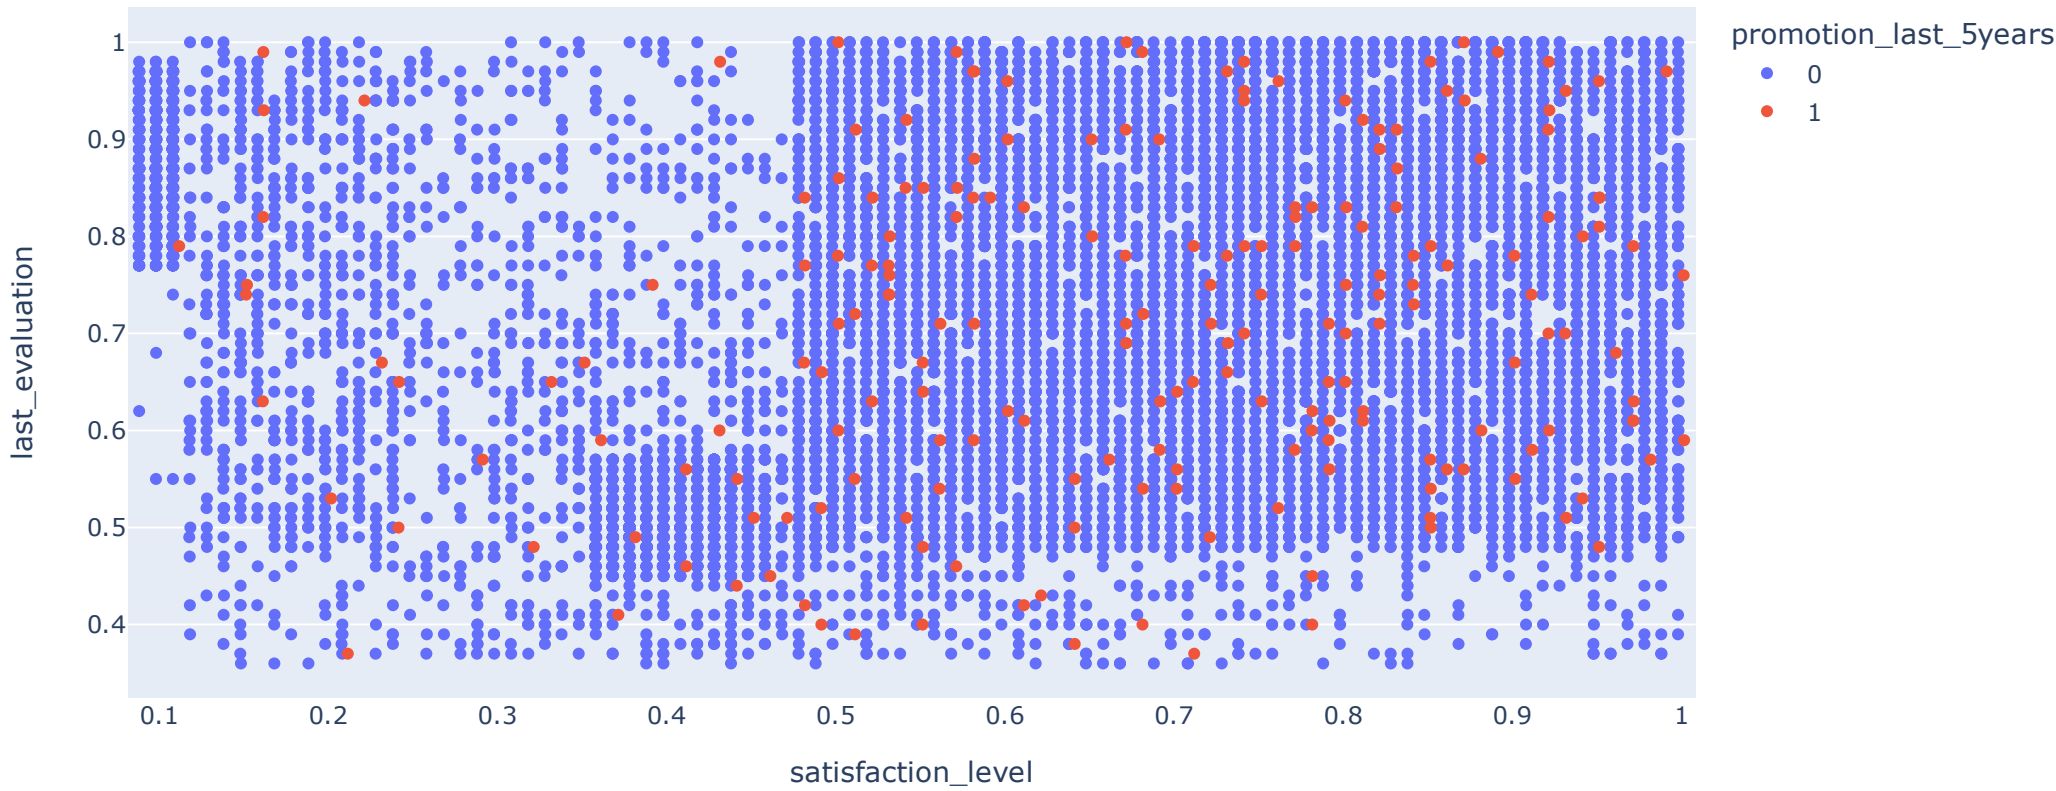
promotion\_last\_5years and left



[Export to plot.ly »](#)

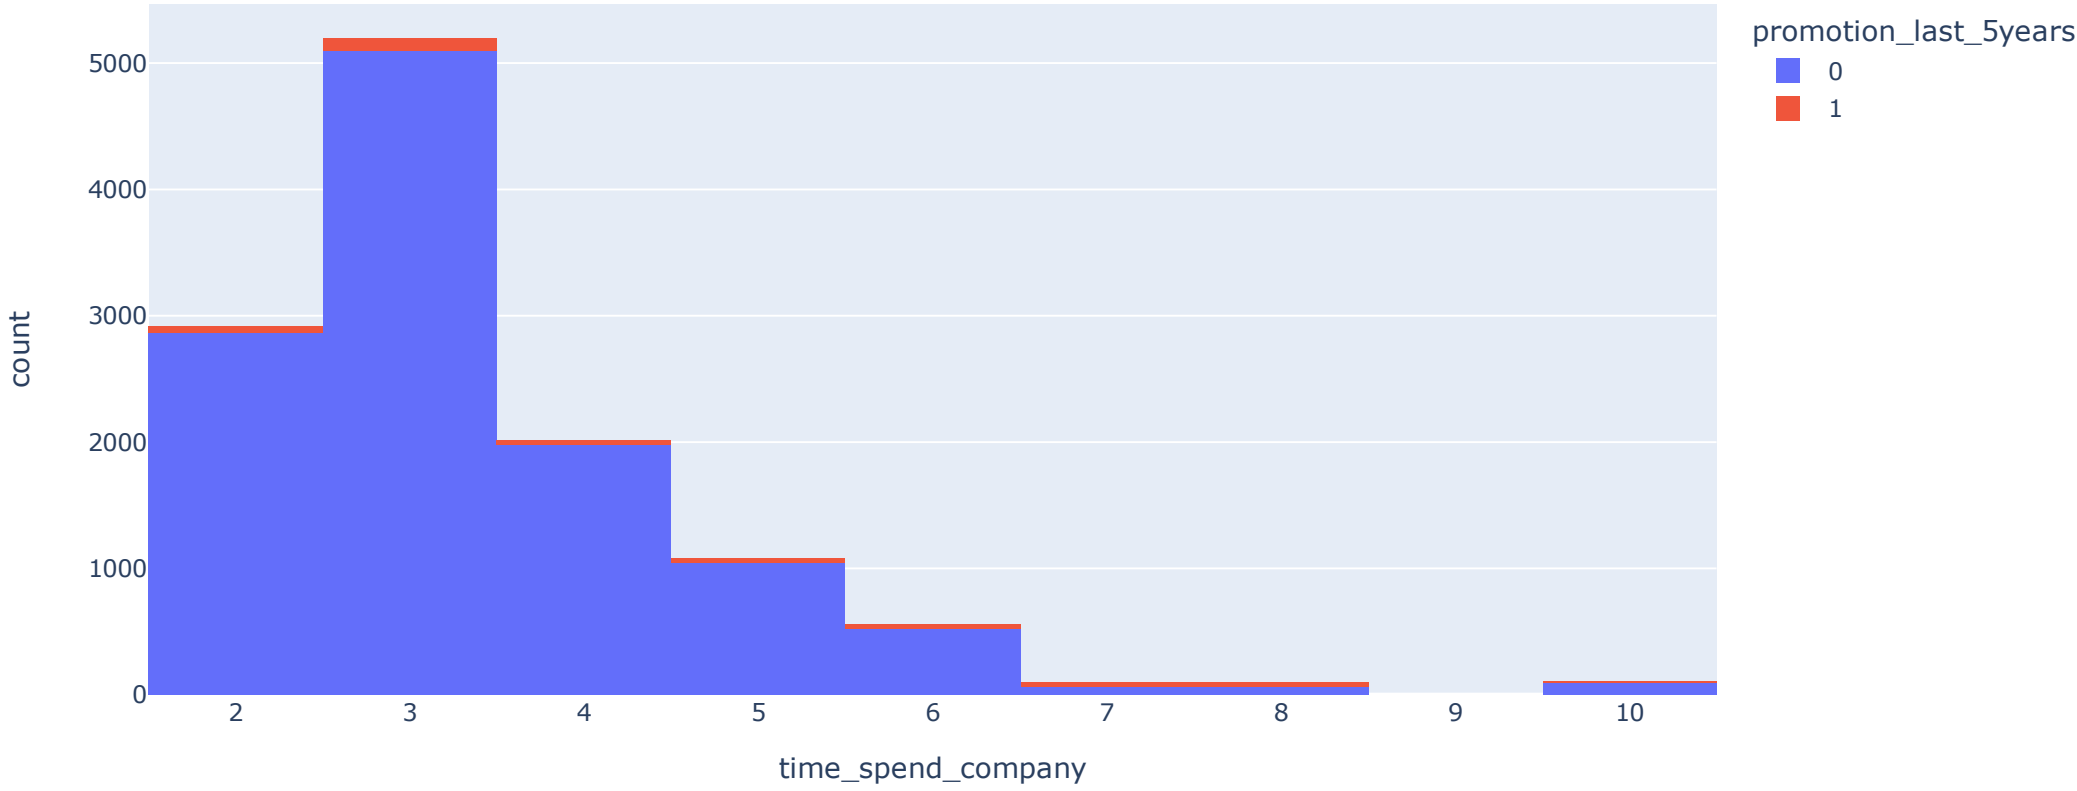
```
In [49]: fig = px.strip(df, x = 'satisfaction_level', y = 'last_evaluation', color = 'promotion_last_5years',
                    title = "'satisfaction_level' & 'last_evaluation'")
fig.show()
```

'satisfaction\_level' & 'last\_evaluation'




```
In [50]: px.histogram(df, x = df['time_spend_company'], color='promotion_last_5years', title = 'time_spend_company')
```

time\_spend\_company



```
In [51]: px.histogram(df[df['promotion_last_5years'] == 1], x = df['time_spend_company'], title = 'promotion_last_5years')
```

promotion\_last\_5years



- 'promotion\_last\_5years' column has binary type values.
- Receiving a promotion in the last 5 working years is not determinative in terms of leaving or continuing to work.


'department' Column

```
In [52]: cprint("Have a First Look to 'department' Column",'green', 'on_black')
first_look('department')

Have a First Look to 'department' Column
column name      : department
-----
Per_of_Nulls    : % 0.0
Num_of_Nulls    : 0
Num_of_Uniques  : 10
Duplicates      : 0
sales           : 3239
technical       : 2244
support         : 1821
IT              : 976
RandD           : 694
product_mng     : 686
marketing       : 673
accounting      : 621
hr              : 601
management      : 436
Name: department, dtype: int64

In [53]: fig = px.pie(df, values = df['department'].value_counts(),
                      names = (df['department'].value_counts()).index,
                      title = '"department" Column Distribution')
fig.show()
```

"department" Column Distribution



```
In [54]: cprint('left, not_left values and left percentage','green', 'on_red')
df_dep = pd.DataFrame(pd.crosstab(df['department'], df['left']))
df_dep.rename(columns = {0 : 'not_left', 1 : 'left'}, inplace = True)
df_dep = df_dep.assign(total = lambda x: (x['not_left'] + x['left']))
df_dep = df_dep.assign(left_percentage = lambda x: (x['left'] / x['total'] * 100))
df_dep


left, not_left values and left percentage

Out[54]:
```

department	left	not_left	left	total	left_percentage
IT	818	158	976	16.189	
RandD	609	85	694	12.248	
accounting	512	109	621	17.552	
hr	488	113	601	18.802	
management	384	52	436	11.927	
marketing	561	112	673	16.642	
product_mng	576	110	686	16.035	
sales	2689	550	3239	16.981	
support	1509	312	1821	17.133	
technical	1854	390	2244	17.380	

```
In [55]: pd.crosstab(df['department'], df['left']).iplot(kind='bar', title = 'department and left')
```

department and left



[Export to plot.ly »](#)

- It is not observed that the departments worked alone have an effect on the left decision.
- It is seen that the left percentages of the departments are similar.


'salary' Column

```
In [56]: cprint("Have a First Look to 'salary' Column",'green', 'on_black')
first_look('salary')

Have a First Look to 'salary' Column
column name      : salary
-----
Per_of_Nulls    : % 0.0
Num_of_Nulls    : 0
Num_of_Uniques  : 3
Duplicates      : 0
low             : 5740
medium          : 5261
high            : 990
Name: salary, dtype: int64


In [57]: fig = px.pie(df, values = df['salary'].value_counts(),
names = (df['salary'].value_counts()).index,
title = "'salary' Column Distribution")
fig.show()
```

"salary" Column Distribution



```
In [58]: pd.crosstab(df['salary'], df['left']).iplot(kind='bar', title = 'salary and left')
```

salary and left



[Export to plot.ly »](#)

- It is seen that the left percentages of the salary are similar.
- Even if it is small, there is an increase in the form of high-medium-low according to the salary status.

**\*Let's go on with the examination of numerical and categorical columns.\***



```
In [59]: numerical= df.drop(['left'], axis = 1).select_dtypes('number').columns

categorical = df.select_dtypes('object').columns

print('-----')
print(f'Numerical Columns: {df[numerical].columns}')
print(f'Categorical Columns: {df[categorical].columns}')
print('-----')

-----
Numerical Columns: Index(['satisfaction_level', 'last_evaluation', 'number_project',
                          'average_montly_hours', 'time_spend_company', 'work_accident',
                          'promotion_last_5years'],
                          dtype='object')
Categorical Columns: Index(['department', 'salary'], dtype='object')
-----


In [60]: cprint("The describe values of the numerical columns",'green', 'on_black')
df[numerical].describe().T.style.background_gradient(subset = ['mean','std','50%','count'], cmap = 'RdPu')
```

The describe values of the numerical columns

	count	mean	std	min	25%	50%	75%	max
satisfaction_level	11991.000000	0.629658	0.241070	0.090000	0.480000	0.660000	0.820000	1.000000
last_evaluation	11991.000000	0.716683	0.168343	0.360000	0.570000	0.720000	0.860000	1.000000
number_project	11991.000000	3.802852	1.163238	2.000000	3.000000	4.000000	5.000000	7.000000
average_montly_hours	11991.000000	200.473522	48.727813	96.000000	157.000000	200.000000	243.000000	310.000000
time_spend_company	11991.000000	3.364857	1.330240	2.000000	3.000000	3.000000	4.000000	10.000000
work_accident	11991.000000	0.154282	0.361234	0.000000	0.000000	0.000000	0.000000	1.000000
promotion_last_5years	11991.000000	0.016929	0.129012	0.000000	0.000000	0.000000	0.000000	1.000000

```
In [61]: df[numerical].iplot(kind = 'histogram', subplots = True, bins = 50, title = 'Histogram visualization of the numerical columns')
```


Histogram visualization of the numerical columns



[Export to plot.ly »](#)

```
In [62]: for i in numerical:
df[i].iplot(kind = 'box', title = i, boxpoints = 'all')
```


satisfaction\_level



[Export to plot.ly »](#)




last\_evaluation




[Export to plot.ly »](#)

number\_project




[Export to plot.ly »](#)

average\_monthly\_hours




[Export to plot.ly »](#)

time\_spend\_company



[Export to plot.ly »](#)

work\_accident



[Export to plot.ly »](#)


promotion\_last\_5years



[Export to plot.ly »](#)

```
In [63]: cprint("The pairplot visualization of the numerical columns",'green', 'on_red')
sns.pairplot(df, hue = "left", corner = True);
```


The pairplot visualization of the numerical columns



```
In [64]: cprint("Heatmap of the numerical columns",'green', 'on_red')

plt.figure(figsize = (12, 8))
sns.heatmap (df.corr(), annot = True, fmt = '.2f', vmin = -1, vmax = 1)
plt.xticks(rotation = 45);
```

Heatmap of the numerical columns



```
In [65]: cprint("Multicollinearity among the features",'green', 'on_red')


df_temp = df.corr()

count = 'Done'
feature = []
collinear= []
for col in df_temp.columns:
    for i in df_temp.index:
        if (df_temp[col][i] > .9 and df_temp[col][i] < 1) or (df_temp[col][i] < -.9 and df_temp[col][i] > -1) :
            feature.append(col)
            collinear.append(i)
            print(Fore.RED + f'\033[1mmulticollinearity alert in between\033[0m {col} - {i}')
        else:
            print(f'For {col} and {i}, there is NO multicollinearity problem')

print('\033[1mThe number of strong corelated features:\033[0m', count)
```

Multicollinearity among the features  
For satisfaction\_level and satisfaction\_level, there is NO multicollinearity problem  
For satisfaction\_level and last\_evaluation, there is NO multicollinearity problem  
For satisfaction\_level and number\_project, there is NO multicollinearity problem  
For satisfaction\_level and average\_monthly\_hours, there is NO multicollinearity problem  
For satisfaction\_level and time\_spend\_company, there is NO multicollinearity problem  
For satisfaction\_level and work\_accident, there is NO multicollinearity problem  
For satisfaction\_level and promotion\_last\_5years, there is NO multicollinearity problem  
For satisfaction\_level and left, there is NO multicollinearity problem  
For last\_evaluation and satisfaction\_level, there is NO multicollinearity problem  
For last\_evaluation and last\_evaluation, there is NO multicollinearity problem  
For last\_evaluation and number\_project, there is NO multicollinearity problem  
For last\_evaluation and average\_monthly\_hours, there is NO multicollinearity problem  
For last\_evaluation and time\_spend\_company, there is NO multicollinearity problem  
For last\_evaluation and work\_accident, there is NO multicollinearity problem  
For last\_evaluation and promotion\_last\_5years, there is NO multicollinearity problem  
For last\_evaluation and left, there is NO multicollinearity problem  
For number\_project and satisfaction\_level, there is NO multicollinearity problem  
For number\_project and last\_evaluation, there is NO multicollinearity problem  
For number\_project and number\_project, there is NO multicollinearity problem  
For number\_project and average\_monthly\_hours, there is NO multicollinearity problem  
For number\_project and time\_spend\_company, there is NO multicollinearity problem  
For number\_project and work\_accident, there is NO multicollinearity problem  
For number\_project and promotion\_last\_5years, there is NO multicollinearity problem  
For number\_project and left, there is NO multicollinearity problem  
For average\_monthly\_hours and satisfaction\_level, there is NO multicollinearity problem  
For average\_monthly\_hours and last\_evaluation, there is NO multicollinearity problem  
For average\_monthly\_hours and number\_project, there is NO multicollinearity problem  
For average\_monthly\_hours and average\_monthly\_hours, there is NO multicollinearity problem  
For average\_monthly\_hours and time\_spend\_company, there is NO multicollinearity problem  
For average\_monthly\_hours and work\_accident, there is NO multicollinearity problem  
For average\_monthly\_hours and promotion\_last\_5years, there is NO multicollinearity problem  
For average\_monthly\_hours and left, there is NO multicollinearity problem  
For time\_spend\_company and satisfaction\_level, there is NO multicollinearity problem  
For time\_spend\_company and last\_evaluation, there is NO multicollinearity problem  
For time\_spend\_company and number\_project, there is NO multicollinearity problem  
For time\_spend\_company and average\_monthly\_hours, there is NO multicollinearity problem  
For time\_spend\_company and time\_spend\_company, there is NO multicollinearity problem  
For time\_spend\_company and work\_accident, there is NO multicollinearity problem  
For time\_spend\_company and promotion\_last\_5years, there is NO multicollinearity problem  
For time\_spend\_company and left, there is NO multicollinearity problem  
For work\_accident and satisfaction\_level, there is NO multicollinearity problem  
For work\_accident and last\_evaluation, there is NO multicollinearity problem  
For work\_accident and number\_project, there is NO multicollinearity problem  
For work\_accident and average\_monthly\_hours, there is NO multicollinearity problem  
For work\_accident and time\_spend\_company, there is NO multicollinearity problem  
For work\_accident and work\_accident, there is NO multicollinearity problem  
For work\_accident and promotion\_last\_5years, there is NO multicollinearity problem  
For work\_accident and left, there is NO multicollinearity problem  
For promotion\_last\_5years and satisfaction\_level, there is NO multicollinearity problem  
For promotion\_last\_5years and last\_evaluation, there is NO multicollinearity problem  
For promotion\_last\_5years and number\_project, there is NO multicollinearity problem  
For promotion\_last\_5years and average\_monthly\_hours, there is NO multicollinearity problem  
For promotion\_last\_5years and time\_spend\_company, there is NO multicollinearity problem  
For promotion\_last\_5years and work\_accident, there is NO multicollinearity problem  
For promotion\_last\_5years and promotion\_last\_5years, there is NO multicollinearity problem  
For promotion\_last\_5years and left, there is NO multicollinearity problem  
For left and satisfaction\_level, there is NO multicollinearity problem  
For left and last\_evaluation, there is NO multicollinearity problem  
For left and number\_project, there is NO multicollinearity problem  
For left and average\_monthly\_hours, there is NO multicollinearity problem  
For left and time\_spend\_company, there is NO multicollinearity problem  
For left and work\_accident, there is NO multicollinearity problem  
For left and promotion\_last\_5years, there is NO multicollinearity problem  
For left and left, there is NO multicollinearity problem  
The number of strong corelated features: Done

```
In [66]: df.corr()['left'].sort_values().drop('left').iplot(kind = 'barh');
```



[Export to plot.ly »](#)

**\*Based on the examinations made above,\***

- There is no multicollinearity problem among the features.
- We have weak level correlation between the numerical features and the target column.
- Also there is weak level correlation between the columns.
- Target variable demonstrates a slight negative correlation with the variables of "promotion\_last\_5years", "work\_accident", "satisfaction\_level",
- Target variable demonstrates slight positive correlation with the variables of 'time\_spend\_company', 'average\_monthly\_hours', 'number\_project' and 'last\_evaluation'.
- satisfaction\_level has more influence on the decision to leave the work than the other columns.

# 5 - DATA VISUALIZATION

- Employees Left
- Determine Number of Projects
- Determine Time Spent in Company


## - Subplots of Features

We can search for answers to the following questions using data visualization methods. Based on these responses, we can develop comments about the factors that cause churn.

- How does the promotion status affect employee churn?
- How does years of experience affect employee churn?
- How does workload affect employee churn?
- How does the salary level affect employee churn?

### 5.1 - Employees Left

Let's check how many employees were left?  
Here, we can plot a bar graph using Matplotlib. The bar graph is suitable for showing discrete variable counts.




### 5.2 - Number of Projects

Similarly, we can also plot a bar graph to count the number of employees deployed on how many projects?







## 5.3 - Time Spent in the Company


Similarly, we can also plot a bar graph to count the number of employees have based on how much experience?

```
In [71]: cprint('"time_spend_company" Column Distribution','green', 'on_red')
df.time_spend_company.value_counts()

"time_spend_company" Column Distribution
Out[71]: 3    5190
         2    2910
         4    2005
         5    1062
         6     542
        10     107
         7      94
         8      81
Name: time_spend_company, dtype: int64

In [72]: cprint('"time_spend_company" Column Distribution','green', 'on_red')
fig = plt.figure(figsize = (11,6))
ax = fig.add_axes([0,0,1,1])
df.time_spend_company.value_counts().plot(kind = "bar", color = "pink")
plt.title('"time_spend_company" Column Distribution')
plt.xlabel('time_spend_company')
plt.ylabel('Number of Employees')
plt.xticks(rotation = 0)
for index,value in enumerate(df.time_spend_company.value_counts().sort_values(ascending=False)):
    plt.text(index, value, f'{value}', ha = 'center', va = 'bottom', fontsize = 13)
plt.show()

"time_spend_company" Column Distribution
```




## 5.4 - Subplots of Features

We can use the methods of the plotly.


```
In [73]: for i in df:
         df[i].iplot(kind = 'histogram', subplots = True, bins = 50, title = 'Subplots of Features')
```

Subplots of Features




[Export to plot.ly »](#)

Subplots of Features




[Export to plot.ly »](#)

Subplots of Features




[Export to plot.ly »](#)

Subplots of Features




[Export to plot.ly »](#)

Subplots of Features




[Export to plot.ly »](#)

Subplots of Features




[Export to plot.ly »](#)

Subplots of Features



[Export to plot.ly »](#)


Subplots of Features



[Export to plot.ly »](#)




Subplots of Features



[Export to plot.ly »](#)

Subplots of Features



[Export to plot.ly »](#)

```
In [102]: ## Still Data visualization had to be good
```

## 6 - DATA PRE-PROCESSING

- Label Encoding
- Scaling

### 6.1 - Label Encoding

Lots of machine learning algorithms require numerical input data, so you need to represent categorical columns in a numerical column. In order to encode this data, you could map each value to a number. e.g. Salary column's value can be represented as low:0, medium:1, and high:2. This process is known as label encoding, and sklearn conveniently will do this for you using LabelEncoder.

```
In [74]: cprint('New df for Kmeans clustering','green', 'on_black')
df1 = df.drop('left', axis = 1)
df1.head(1)
```

New df for Kmeans clustering

	satisfaction_level	last_evaluation	number_project	average_montly_hours	time_spend_company	work_accident	promotion_last_5years	department	salary
0	0.380	0.530	2	157	3	0	0	sales	low

```
In [75]: cprint('New df after getting dummied','green', 'on_black')
df1 = pd.get_dummies(df1, columns = ['department','salary'], drop_first = True)
df1.head(1)
```

New df after getting dummied

	satisfaction_level	last_evaluation	number_project	average_montly_hours	time_spend_company	work_accident	promotion_last_5years	department_RandD	department_accounting	department_sales	low	medium	high
0	0.380	0.530	2	157	3	0	0	0	0	0	1	0	0

### 6.2 - Scalling

Some machine learning algorithms are sensitive to feature scaling while others are virtually invariant to it. Machine learning algorithms like linear regression, logistic regression, neural network, etc. that use gradient descent as an optimization technique require data to be scaled. Also distance algorithms like KNN, K-means, and SVM are most affected by the range of features. This is because behind the scenes they are using distances between data points to determine their similarity.

Scaling Types:

- Normalization: Normalization is a scaling technique in which values are shifted and rescaled so that they end up ranging between 0 and 1. It is also known as Min-Max scaling.
- Standardization: Standardization is another scaling technique where the values are centered around the mean with a unit standard deviation. This means that the mean of the attribute becomes zero and the resultant distribution has a unit standard deviation.

Click here for more on scaling. (<https://towardsdatascience.com/all-about-feature-scaling-bcc0ad75cb35>)

```
In [76]: df1.head()
```

Out[76]:

	satisfaction_level	last_evaluation	number_project	average_monthly_hours	time_spend_company	work_accident	promotion_last_5years	department_RandD	department_accounting
0	0.380	0.530	2	157	3	0	0	0	(
1	0.800	0.860	5	262	6	0	0	0	(
2	0.110	0.880	7	272	4	0	0	0	(
3	0.720	0.870	5	223	5	0	0	0	(
4	0.370	0.520	2	159	3	0	0	0	(

In [77]:

```
cprint('Scaling','green', 'on_black')
scaler = MinMaxScaler()
scaler.fit(df1)
#Store it separately for clustering
df1_scaled= scaler.transform(df1)
```

Scaling

In [78]:

```
df1_scaled
```

Out[78]:

```
array([[0.31868132, 0.265625 , 0.        , ..., 0.        , 1.        ,
        0.        ],
       [0.78021978, 0.78125 , 0.6        , ..., 0.        , 0.        ,
        1.        ],
       [0.02197802, 0.8125 , 1.        , ..., 0.        , 0.        ,
        1.        ],
       ...,
       [0.83516484, 0.28125 , 0.2        , ..., 0.        , 0.        ,
        0.        ],
       [0.26373626, 0.453125 , 0.2        , ..., 0.        , 0.        ,
        0.        ],
       [0.45054945, 0.578125 , 0.4        , ..., 0.        , 1.        ,
        0.        ]])
```

## 7 - CLUSTER ANALYSIS

- Find the optimal number of clusters (k) using the elbow method for for K-means.
- Determine the clusters by using K-Means then Evaluate predicted results.

Cluster analysis or clustering is the task of grouping a set of objects in such a way that objects in the same group (called a cluster) are more similar (in some sense) to each other than to those in other groups (clusters). It is a main task of exploratory data mining, and a common technique for statistical data analysis, used in many fields, including pattern recognition, image analysis, information retrieval, bioinformatics, data compression, computer graphics and machine learning.

Cluster Analysis ([https://en.wikipedia.org/wiki/Cluster\\_analysis](https://en.wikipedia.org/wiki/Cluster_analysis))

Cluster Analysis2 (<https://realpython.com/k-means-clustering-python/>)

### The Elbow Method

- "Elbow Method" can be used to find the optimum number of clusters in cluster analysis. The elbow method is used to determine the optimal number of clusters in k-means clustering. The elbow method plots the value of the cost function produced by different values of k. If k increases, average distortion will decrease, each cluster will have fewer constituent instances, and the instances will be closer to their respective centroids. However, the improvements in average distortion will decline as k increases. The value of k at which improvement in distortion declines the most is called the elbow, at which we should stop dividing the data into further clusters.

[The Elbow Method]([https://en.wikipedia.org/wiki/Elbow\\_method\\_\(clustering\)](https://en.wikipedia.org/wiki/Elbow_method_(clustering)) ([https://en.wikipedia.org/wiki/Elbow\\_method\\_\(clustering\)](https://en.wikipedia.org/wiki/Elbow_method_(clustering)))

The Elbow Method2 (<https://medium.com/@mudgalvivek2911/machine-learning-clustering-elbow-method-4e8c2b404a5d>)


KMeans (<https://towardsdatascience.com/understanding-k-means-clustering-in-machine-learning-6a6e67336aa1>)

Let's find out the groups of employees who left. You can observe that the most important factor for any employee to stay or leave is satisfaction and performance in the company. So let's bunch them in the group of people using cluster analysis.

In [79]:


```
#First : Get the Best KMeans
ks = range(1,10)
inertias=[]
for k in ks :
    # Create a KMeans clusters
    kc = KMeans(n_clusters=k,random_state=1)
    kc.fit(df1_scaled)
    inertias.append(kc.inertia_)

# Plot ks vs inertias
f, ax = plt.subplots(figsize=(8, 6))
plt.plot(ks, inertias, '-o')
plt.xlabel('Number of clusters, k')
plt.ylabel('Inertia')
plt.xticks(ks)
plt.style.use('ggplot')
plt.title('What is the Best Number for KMeans ?')
plt.show()
```



In [80]:

```
from yellowbrick.cluster import KElbowVisualizer
kmeans = KMeans()
visu = KElbowVisualizer(kmeans, k = (1,10))
visu.fit(df1_scaled)
visu.show();
```



```
In [81]: cprint("Silhouette Scores",'green', 'on_red')

ssd =[]

K = range(2,10)

for k in K:
    model = KMeans(n_clusters=k)
    model.fit(df1_scaled)
    ssd.append(model.inertia_)
    print(f'Silhouette Score for {k} clusters: {silhouette_score(df1_scaled, model.labels_)})')
```


```
Silhouette Scores
Silhouette Score for 2 clusters: 0.24293849820807237
Silhouette Score for 3 clusters: 0.19250986825854602
Silhouette Score for 4 clusters: 0.18041466053129487
Silhouette Score for 5 clusters: 0.21174582903978412
Silhouette Score for 6 clusters: 0.2405335474200126
Silhouette Score for 7 clusters: 0.2686037509709165
Silhouette Score for 8 clusters: 0.29285588884780916
Silhouette Score for 9 clusters: 0.3069720256203487
```

```
In [83]: cprint("Silhouette Plot for K=3",'green', 'on_red')

from sklearn.cluster import KMeans
from yellowbrick.cluster import SilhouetteVisualizer

model_3 = KMeans(n_clusters = 3, random_state = 101)
visualizer = SilhouetteVisualizer(model_3)
visualizer.fit(df1_scaled)
visualizer.poof();
```

Silhouette Plot for K=3




```
In [84]: cprint("Silhouette Plot for K=4",'green', 'on_red')

from sklearn.cluster import KMeans
from yellowbrick.cluster import SilhouetteVisualizer

model_4 = KMeans(n_clusters = 4, random_state = 101)
visualizer = SilhouetteVisualizer(model_4)
visualizer.fit(df1_scaled)
visualizer.poof();
```

Silhouette Plot for K=4



According to the silhouette score, clustering according to the K=2, K=3 and K=4 are seen above.

- For K=3, 0 labelled cluster is below the average silhouette score.
- For K=4, 0 and 3 labelled clusters are below the average silhouette score.
- For K=3 (According to Elbow) and for K=4 (According to the silhouette score) clustering is not suitable for our dataset.


Let's see How it is when K=2 (According to our target variable classes)

In [85]: `cprint("Silhouette Plot for K=2",'green', 'on_red')`

```
from sklearn.cluster import KMeans
from yellowbrick.cluster import SilhouetteVisualizer

model_2 = KMeans(n_clusters = 2, random_state = 101)
visualizer = SilhouetteVisualizer(model_2)
visualizer.fit(df1_scaled)
visualizer.poof();
```

Silhouette Plot for K=2



In [86]: `cprint("KMeans Clustering with K=2",'green', 'on_red')`

```
k_means_model2 = KMeans(n_clusters = 2, random_state = 101)
k_means_model2.fit_predict(df1_scaled)
labels = k_means_model2.labels_
labels
```

KMeans Clustering with K=2

Out[86]: `array([1, 0, 0, ..., 0, 0, 1])`

In [87]: `cprint("Predicted clusters on our dataframe",'green', 'on_red')`

```
df['predicted_clusters'] = labels
df
```

Predicted clusters on our dataframe

Out[87]:

	satisfaction_level	last_evaluation	number_project	average_monthly_hours	time_spend_company	work_accident	promotion_last_5years	department	salary	left	predicted
0	0.380	0.530	2	157	3	0	0	sales	low	1	
1	0.800	0.860	5	262	6	0	0	sales	medium	1	
2	0.110	0.880	7	272	4	0	0	sales	medium	1	
3	0.720	0.870	5	223	5	0	0	sales	low	1	
4	0.370	0.520	2	159	3	0	0	sales	low	1	
...	...	...	...	...	...	...	...	...	...	...	
11995	0.900	0.550	3	259	10	1	1	management	high	0	
11996	0.740	0.950	5	266	10	0	1	management	high	0	
11997	0.850	0.540	3	185	10	0	1	management	high	0	
11998	0.330	0.650	3	172	10	0	1	marketing	high	0	
11999	0.500	0.730	4	180	3	0	0	IT	low	0	

11991 rows × 11 columns


In [88]: `cprint('"predicted_clusters" value counts','green', 'on_red')`  
`df['predicted_clusters'].value_counts()`

"predicted\_clusters" value counts

Out[88]: `0 6251`  
`1 5740`  
`Name: predicted_clusters, dtype: int64`


```
In [89]: fig = px.pie(df, values = df['predicted_clusters'].value_counts(),
names = (df['predicted_clusters'].value_counts().index,
title = 'Predicted_Clusters Distribution')
fig.show()
```

Predicted\_Clusters Distribution




```
In [90]: fig = px.pie(df, values = df[df['left']==0]['predicted_clusters'].value_counts(),
names = df[df['left']==0]['predicted_clusters'].value_counts().index,
title = 'Predicted_Clusters & left==0 Distribution')
fig.show()
```

Predicted\_Clusters & left==0 Distribution



```
In [91]: fig = px.pie(df, values = df[df['left']==1]['predicted_clusters'].value_counts(),
names = df[df['left']==1]['predicted_clusters'].value_counts().index,
title = 'Predicted_Clusters & left==1 Distribution')
fig.show()
```

Predicted\_Clusters & left==1 Distribution



```
In [92]: cprint('Mean values according to the left','green', 'on_red')
df.groupby('left').mean()
```

Mean values according to the left

Out[92]:	satisfaction_level	last_evaluation	number_project	average_monthly_hours	time_spend_company	work_accident	promotion_last_5years	predicted_clusters	
	left								
	0	0.667	0.716	3.787	198.943	3.262	0.174	0.019	0.457
	1	0.440	0.722	3.883	208.162	3.881	0.053	0.004	0.590

```
In [93]: cprint('Mean values according to the left and predicted clusters','green', 'on_red')
df.groupby(['left', 'predicted_clusters']).mean()
```

Mean values according to the left and predicted clusters

Out[93]:

		satisfaction_level	last_evaluation	number_project	average_monthly_hours	time_spend_company	work_accident	promotion_last_5years	
left	predicted_clusters								
	0	0	0.666	0.714	3.788	199.173	3.334	0.168	0.030
1		1	0.669	0.718	3.786	198.669	3.176	0.182	0.007
		0	0.440	0.722	3.911	209.366	3.886	0.059	0.004
		1	0.441	0.722	3.865	207.325	3.878	0.049	0.004

In [94]:

```
pd.crosstab(df['left'], df['predicted_clusters']).iplot(kind="bar", title = 'Compare (left vs predicted_clusters)',
xTitle = 'left & clusters', yTitle = 'counts')
```



[Export to plot.ly »](#)

In [95]:

```
cprint('Mean values according to the predicted clusters','green', 'on_red')
df.groupby('predicted_clusters').mean()
```

Mean values according to the predicted clusters

Out[95]:

	satisfaction_level	last_evaluation	number_project	average_monthly_hours	time_spend_company	work_accident	promotion_last_5years	left
predicted_clusters								
0	0.636	0.715	3.804	200.505	3.406	0.154	0.026	0.131
1	0.623	0.719	3.802	200.439	3.320	0.155	0.007	0.205

In [96]:

```
cprint('Mean values according to the predicted clusters and left','green', 'on_red')
df.groupby(['predicted_clusters', 'left']).mean()
```


Mean values according to the predicted clusters and left

Out[96]:

		satisfaction_level	last_evaluation	number_project	average_monthly_hours	time_spend_company	work_accident	promotion_last_5years	
predicted_clusters	left								
	0	0	0.666	0.714	3.788	199.173	3.334	0.168	0.030
1		1	0.440	0.722	3.911	209.366	3.886	0.059	0.004
		0	0.669	0.718	3.786	198.669	3.176	0.182	0.007
		1	0.441	0.722	3.865	207.325	3.878	0.049	0.004

In [97]:

```
pd.crosstab(df['predicted_clusters'],
df['left']).iplot(kind="bar", title = 'Compare (predicted_clusters vs left)',
xTitle = 'clusters & left', yTitle = 'counts')
```



[Export to plot.ly »](#)

In [98]:

```
cprint('Mean values of the features according to the predicted clusters','green', 'on_red')


sns.lineplot(data = df.iloc[:, [0, 1, 2, 3, 4, 5, 6, 7, 8, 10]].groupby("predicted_clusters").mean().T)
plt.xticks(rotation = 60)
```

Mean values of the features according to the predicted clusters

Out[98]:

[0, 1, 2, 3, 4, 5, 6],
[Text(0, 0, 'satisfaction_level'),
Text(1, 0, 'last_evaluation'),
Text(2, 0, 'number_project'),
Text(3, 0, 'average_monthly_hours'),
Text(4, 0, 'time_spend_company'),
Text(5, 0, 'work_accident'),
Text(6, 0, 'promotion_last_5years')]]





As seen above, the columns in the data set do not separate from each other. All columns are intertwined with each other.

As seen above it is visually obvious that **clustering is not a good approach to our dataset**.

From now on we are going to use **classification models to make churn predictions**.

## 8 - MODEL BUILDING

- Split Data as Train and Test set
- Built Gradient Boosting Classifier, Evaluate Model Performance and Predict Test Data
- Built K Neighbors Classifier and Evaluate Model Performance and Predict Test Data

### - Built Random Forest Classifier and Evaluate Model Performance and Predict Test Data

Evaluating Model Performance

- Confusion Matrix : You can use scikit-learn metrics module for accuracy calculation. A Confusion matrix is an N x N matrix used for evaluating the performance of a classification model, where N is the number of target classes. The matrix compares the actual target values with those predicted by the machine learning model. This gives us a holistic view of how well our classification model is performing and what kinds of errors it is making.

Confusion Matrix (<https://www.analyticsvidhya.com/blog/2020/04/confusion-matrix-machine-learning/>)

- Yellowbrick: Yellowbrick is a suite of visualization and diagnostic tools that will enable quicker model selection. It's a Python package that combines scikit-learn and matplotlib. Some of the more popular visualization tools include model selection, feature visualization, classification and regression visualization

Yellowbrick (<https://www.analyticsvidhya.com/blog/2018/05/yellowbrick-a-set-of-visualization-tools-to-accelerate-your-model-selection-process/>)

**Here, Dataset is broken into two parts in ratio of 70:30. It means 70% data will used for model training and 30% for model testing.**

```
In [99]: cprint('New df for Classification','green', 'on_red')
df2 = df.drop('predicted_clusters', axis = 1)
df2.head(1)
```

New df for Classification

```
Out[99]:
```

	satisfaction_level	last_evaluation	number_project	average_monthly_hours	time_spend_company	work_accident	promotion_last_5years	department	salary	left
0	0.380	0.530	2	157	3	0	0	sales	low	1

```
In [100]: cprint('New df after getting dummied','green', 'on_red')
df2 = pd.get_dummies(df2, columns = ['department','salary'], drop_first = True)
df2.head(1)
```

New df after getting dummied

```
Out[100]:
```

	satisfaction_level	last_evaluation	number_project	average_monthly_hours	time_spend_company	work_accident	promotion_last_5years	left	department_RandD	department_acc
0	0.380	0.530	2	157	3	0	0	1		0

### 8.1 - Splitting Data as Train & Test

```
In [101]: X = df2.drop('left', axis = 1)
y = df2['left']
```

```
In [102]: X_train, X_test, y_train, y_test = train_test_split(X, y, stratify = y, test_size = 0.3, random_state = 101)
```

```
In [103]: scaler = MinMaxScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```
In [104]: def eval(model, X_train, X_test):
    y_pred = model.predict(X_test)
    y_pred_train = model.predict(X_train)

    print(confusion_matrix(y_test, y_pred))
    print("Test_Set")
    print(classification_report(y_test,y_pred))
    print("Train_Set")
    print(classification_report(y_train,y_pred_train))
```

```
In [105... def train_val(y_train, y_train_pred, y_test, y_pred):

    scores = {"train_set": {"Accuracy" : accuracy_score(y_train, y_train_pred),
                          "Precision" : precision_score(y_train, y_train_pred),
                          "Recall" : recall_score(y_train, y_train_pred),
                          "f1" : f1_score(y_train, y_train_pred)},

              "test_set": {"Accuracy" : accuracy_score(y_test, y_pred),
                          "Precision" : precision_score(y_test, y_pred),
                          "Recall" : recall_score(y_test, y_pred),
                          "f1" : f1_score(y_test, y_pred)}}

    return pd.DataFrame(scores)
```

8.2 - Gradient Boosting Classifier

8.2.1 Model Building

```
In [106... GB_model = GradientBoostingClassifier(random_state = 101)
GB_model.fit(X_train, y_train)
y_pred = GB_model.predict(X_test)
y_train_pred = GB_model.predict(X_train)

GB_model_f1 = f1_score(y_test, y_pred)
GB_model_acc = accuracy_score(y_test, y_pred)
GB_model_recall = recall_score(y_test, y_pred)
GB_model_auc = roc_auc_score(y_test, y_pred)
```

8.2.2 Evaluating Model Performance

```
In [107... print("GB_Model")
print ("-----")
eval(GB_model, X_train, X_test)
```

GB\_Model

-----

[[2974 27]
[ 45 552]]

Test\_Set

	precision	recall	f1-score	support
0	0.99	0.99	0.99	3001
1	0.95	0.92	0.94	597
accuracy			0.98	3598
macro avg	0.97	0.96	0.96	3598
weighted avg	0.98	0.98	0.98	3598

Train\_Set


	precision	recall	f1-score	support
0	0.99	0.99	0.99	6999
1	0.97	0.93	0.95	1394
accuracy			0.98	8393
macro avg	0.98	0.96	0.97	8393
weighted avg	0.98	0.98	0.98	8393

```
In [108... cprint('GB_model Scores','green', 'on_red')
train_val(y_train, y_train_pred, y_test, y_pred)
```

GB\_model Scores

	train_set	test_set
Accuracy	0.983	0.980
Precision	0.969	0.953
Recall	0.928	0.925
f1	0.948	0.939

```
In [109... from yellowbrick.classifier import ClassPredictionError
visualizer = ClassPredictionError(GB_model)
# Fit the training data to the visualizer
visualizer.fit(X_train, y_train)
# Evaluate the model on the test data
visualizer.score(X_test, y_test)
# Draw visualization
visualizer.poof();
```



8.2.3 Feature Importance for Gradient Boosting Model


```
In [110... GB_feature_imp = pd.DataFrame(index=X.columns, data = GB_model.feature_importances_, columns = ['Importance']).sort_values("Importance", ascending :
GB_feature_imp
```



```
Out[110]:
```

	Importance
satisfaction_level	0.527
number_project	0.138
time_spend_company	0.137
last_evaluation	0.122
average_monthly_hours	0.075
salary_low	0.001
work_accident	0.000
department_management	0.000
department_product_mng	0.000
department_RandD	0.000
department_marketing	0.000
salary_medium	0.000
department_accounting	0.000
department_sales	0.000
department_support	0.000
department_technical	0.000
promotion_last_5years	0.000
department_hr	0.000

```
In [111]: fig = px.bar(GB_feature_imp.sort_values('Importance', ascending = False), x = GB_feature_imp.sort_values('Importance', ascending = False).index, y = 'Importance', title = "Feature Importance", labels = dict(x = "Features", y ="Feature_Importance"))
fig.show()
```



8.2.4 Gradient Boosting Classifier Cross Validation

```
In [112]: GB_cv = GradientBoostingClassifier(random_state = 101)
GB_cv_scores = cross_validate(GB_cv, X_train, y_train,
                              scoring = ['accuracy', 'precision', 'recall', 'f1', 'roc_auc'], cv = 10)
GB_cv_scores = pd.DataFrame(GB_cv_scores, index = range(1, 11))
GB_cv_scores.mean()[2:]

Out[112]: test_accuracy    0.981
test_precision    0.965
test_recall    0.922
test_f1    0.943
test_roc_auc    0.985
dtype: float64
```

8.2.5 Gradient Boosting Classifier GridSearchCV

```
In [113]: param_grid = {"n_estimators": [100, 200, 300],
                        "subsample": [0.5, 1],
                        "max_features" : [None, 2, 3, 4],
                        "learning_rate": [0.001, 0.01, 0.1],
                        'max_depth': [3, 4, 5, 6]}

In [114]: GB_grid = GradientBoostingClassifier(random_state = 101)
GB_grid_model = GridSearchCV(GB_grid, param_grid, scoring = "f1", verbose = 2, n_jobs = -1).fit(X_train, y_train)

Fitting 5 folds for each of 288 candidates, totalling 1440 fits

In [115]: GB_grid_model.best_estimator_

Out[115]: ▼ GradientBoostingClassifier
GradientBoostingClassifier(learning_rate=0.01, max_depth=6, n_estimators=200,
                           random_state=101, subsample=0.5)

In [116]: print(colored('\033[1mBest Parameters of GridSearchCV for Gradient Boosting Model:\033[0m', 'blue'), colored(GB_grid_model.best_params_, 'red'))

Best Parameters of GridSearchCV for Gradient Boosting Model: {'learning_rate': 0.01, 'max_depth': 6, 'max_features': None, 'n_estimators': 200, 'subsample': 0.5}

In [117]: GB_tuned = GradientBoostingClassifier(learning_rate = 0.01,
                                                max_depth = 6,
                                                n_estimators = 200,
                                                subsample = 0.5,
                                                random_state = 101).fit(X_train, y_train)

In [118]: y_pred = GB_tuned.predict(X_test)
y_train_pred = GB_tuned.predict(X_train)

GB_tuned_f1 = f1_score(y_test, y_pred)
GB_tuned_acc = accuracy_score(y_test, y_pred)
GB_tuned_recall = recall_score(y_test, y_pred)
GB_tuned_auc = roc_auc_score(y_test, y_pred)
```

```
In [119... print("GB_tuned")
print ("-----")
eval(GB_tuned, X_train, X_test)

GB_tuned
-----
[[2994    7]
 [  48 549]]
Test_Set
      precision    recall  f1-score   support

      0       0.98       1.00       0.99       3001
      1       0.99       0.92       0.95        597

   accuracy       0.98       0.98       0.98       3598
  macro avg       0.99       0.96       0.97       3598
 weighted avg       0.98       0.98       0.98       3598

Train_Set
      precision    recall  f1-score   support

      0       0.99       1.00       0.99       6999
      1       0.99       0.93       0.96       1394

   accuracy       0.99       0.99       0.99       8393
  macro avg       0.99       0.96       0.97       8393
 weighted avg       0.99       0.99       0.99       8393


In [120... cprint('GB_tuned Scores','green', 'on_red')
train_val(y_train, y_train_pred, y_test, y_pred)
```

GB\_tuned Scores

Out[120]:

	train_set	test_set
Accuracy	0.986	0.985
Precision	0.989	0.987
Recall	0.927	0.920
f1	0.957	0.952

```
In [121... from yellowbrick.classifier import ClassPredictionError
visualizer = ClassPredictionError(GB_tuned)
# Fit the training data to the visualizer
visualizer.fit(X_train, y_train)
# Evaluate the model on the test data
visualizer.score(X_test, y_test)
# Draw visualization
visualizer.poof();
```



8.2.7 Prediction

```
In [122... cprint('GB_tuned Predictions','green', 'on_red')
GB_Pred = {"Actual": y_test, "GB_Pred":y_pred}
GB_Pred = pd.DataFrame.from_dict(GB_Pred)
GB_Pred.head()
```

GB\_tuned Predictions

Out[122]:

	Actual	GB_Pred
3118	0	0
10490	0	0
1106	1	1
3822	0	0
6873	0	0

```
In [123... cprint('Predictions','green', 'on_red')
Model_Preds = GB_Pred
Model_Preds.head()
```

Predictions

Out[123]:

	Actual	GB_Pred
3118	0	0
10490	0	0
1106	1	1
3822	0	0
6873	0	0

8.3 - KNeighbors Classifier

8.3.1 Model Building

```
In [124... KNN_model = KNeighborsClassifier(n_neighbors = 5)
KNN_model.fit(X_train, y_train)
y_pred = KNN_model.predict(X_test)
y_train_pred = KNN_model.predict(X_train)

KNN_model_f1 = f1_score(y_test, y_pred)
KNN_model_acc = accuracy_score(y_test, y_pred)
KNN_model_recall = recall_score(y_test, y_pred)
KNN_model_auc = roc_auc_score(y_test, y_pred)
```

8.3.2 Evaluating Model Performance


```
In [125... print("KNN_Model")
print ("-----")
eval(KNN_model, X_train, X_test)
```

KNN_Model					
-----					
[[2883  118]					
[  83  514]]					
Test_Set					
	precision	recall	f1-score	support	
0	0.97	0.96	0.97	3001	
1	0.81	0.86	0.84	597	
accuracy			0.94	3598	
macro avg	0.89	0.91	0.90	3598	
weighted avg	0.95	0.94	0.94	3598	
Train_Set					
	precision	recall	f1-score	support	
0	0.98	0.98	0.98	6999	
1	0.88	0.88	0.88	1394	
accuracy			0.96	8393	
macro avg	0.93	0.93	0.93	8393	
weighted avg	0.96	0.96	0.96	8393	

```
In [126... cprint('KNN_model Scores','green', 'on_red')
train_val(y_train, y_train_pred, y_test, y_pred)
```

KNN_model Scores			
	train_set	test_set	
Accuracy	0.960	0.944	
Precision	0.879	0.813	
Recall	0.878	0.861	
f1	0.879	0.836	

```
In [127... from yellowbrick.classifier import ClassPredictionError
visualizer = ClassPredictionError(KNN_model)
# Fit the training data to the visualizer
visualizer.fit(X_train, y_train)
# Evaluate the model on the test data
visualizer.score(X_test, y_test)
# Draw visualization
visualizer.poof();
```



8.3.3 KNeighbors Classifier Cross Validation

```
In [128... KNN_cv = KNeighborsClassifier(n_neighbors = 5)
KNN_cv_scores = cross_validate(KNN_cv, X_train, y_train,
                               scoring = ['accuracy', 'precision','recall', 'f1', 'roc_auc'], cv = 10)
KNN_cv_scores = pd.DataFrame(KNN_cv_scores, index = range(1, 11))
KNN_cv_scores.mean()[2:]
```

test_accuracy	0.948
test_precision	0.835
test_recall	0.859
test_f1	0.846
test_roc_auc	0.943
dtype: float64	

8.3.4 Elbow Method for Choosing Reasonable K Values

```
In [129... test_error_rates = []

for k in range(1, 30):
    KNN = KNeighborsClassifier(n_neighbors = k)
    KNN.fit(X_train, y_train)

    y_pred = KNN.predict(X_test)


    test_error = 1 - accuracy_score(y_test, y_pred)
    test_error_rates.append(test_error)

print(test_error_rates)

plt.figure(figsize = (15, 8))
plt.plot(range(1, 30), test_error_rates, color = 'blue', linestyle = '--', marker = 'o',
        markerfacecolor = 'red', markersize = 10)
plt.title('Error Rate vs. K Value')
plt.xlabel('K_values')
plt.ylabel('Error Rate')

plt.hlines(y = 0.04307948860478039, xmin = 0, xmax = 30, colors = 'r', linestyle = "--", label = "K-Value = 2 Line")
plt.hlines(y = 0.04558087826570312, xmin = 0, xmax = 30, colors = 'r', linestyle = "--", label = "K-Value = 4 Line")
plt.hlines(y = 0.055864369093941, xmin = 0, xmax = 30, colors = 'blue', linestyle = "--", label = "Default K-Value = 5 Line")
plt.legend(prop = {"size":14});

[0.053362979433018376, 0.04307948860478039, 0.053085047248471406, 0.04558087826570312, 0.0558643690939411, 0.05141745414118959, 0.0569760978321289
8, 0.05669816564758201, 0.06142301278488049, 0.060867148415786554, 0.06475819899944413, 0.06698165647581988, 0.07031684269038352, 0.071428571428571
4, 0.07559755419677594, 0.07448582545858806, 0.07809894385769867, 0.07865480822679272, 0.08032240133407453, 0.08282379099499726, 0.0839355197331851
4, 0.08643690939410786, 0.0861589772095609, 0.08893829905503059, 0.09032795997776544, 0.09088382434685938, 0.0933852140077821, 0.09310728182323513,
0.09366314619232907]
```



Let's have look to recall values for different K's ranging from 1 to 10

```
In [130... # FIRST A QUICK COMPARISON TO OUR DEFAULT K=5

knn5 = KNeighborsClassifier(n_neighbors = 5)

knn5.fit(X_train,y_train)
pred = knn5.predict(X_test)

print('WITH K=5')
print('-----')
print(confusion_matrix(y_test, pred))
print('-----')
print(classification_report(y_test, pred))

WITH K=5
-----
[[2883  118]
 [  83  514]]
-----
              precision    recall  f1-score   support

      0       0.97        0.96        0.97        3001
      1       0.81        0.86        0.84         597

 accuracy          0.94        0.94        0.94        3598
 macro avg         0.89        0.91        0.90        3598
weighted avg         0.95        0.94        0.94        3598
```

```
In [131... # NOW K=2

knn2 = KNeighborsClassifier(n_neighbors = 2)

knn2.fit(X_train,y_train)
pred = knn2.predict(X_test)

print('WITH K=2')
print('-----')
print(confusion_matrix(y_test, pred))
print('-----')
print(classification_report(y_test, pred))

WITH K=2
-----
[[2951   50]
 [ 105  492]]
-----
              precision    recall  f1-score   support

      0       0.97        0.98        0.97        3001
      1       0.91        0.82        0.86         597

 accuracy          0.96        0.96        0.96        3598
 macro avg         0.94        0.90        0.92        3598
weighted avg         0.96        0.96        0.96        3598
```

In [132...

```
# NOW K=4

knn4 = KNeighborsClassifier(n_neighbors = 4)

knn4.fit(X_train,y_train)
pred = knn4.predict(X_test)

print('WITH K=4')
print('-----')
print(confusion_matrix(y_test, pred))
print('-----')
print(classification_report(y_test, pred))

WITH K=4
-----
[[2929   72]
 [  92 505]]
-----
              precision    recall  f1-score   support

      0       0.97       0.98       0.97       3001
      1       0.88       0.85       0.86       597

 accuracy         0.95
 macro avg       0.92
weighted avg       0.95
```

In [133...

```
# NOW K=6

knn6 = KNeighborsClassifier(n_neighbors = 6)

knn6.fit(X_train,y_train)
pred = knn6.predict(X_test)

print('WITH K=6')
print('-----')
print(confusion_matrix(y_test, pred))
print('-----')
print(classification_report(y_test, pred))

WITH K=6
-----
[[2905   96]
 [  89 508]]
-----
              precision    recall  f1-score   support

      0       0.97       0.97       0.97       3001
      1       0.84       0.85       0.85       597

 accuracy         0.95
 macro avg       0.91
weighted avg       0.95
```

In [134...

```
# NOW K=8

knn8 = KNeighborsClassifier(n_neighbors = 8)

knn8.fit(X_train,y_train)
pred = knn8.predict(X_test)

print('WITH K=8')
print('-----')
print(confusion_matrix(y_test, pred))
print('-----')
print(classification_report(y_test, pred))

WITH K=8
-----
[[2903   98]
 [ 106 491]]
-----
              precision    recall  f1-score   support

      0       0.96       0.97       0.97       3001
      1       0.83       0.82       0.83       597

 accuracy         0.94
 macro avg       0.90
weighted avg       0.94
```

In [135...

```
# NOW K=10

knn10 = KNeighborsClassifier(n_neighbors = 10)

knn10.fit(X_train,y_train)
pred = knn10.predict(X_test)

print('WITH K=10')
print('-----')
print(confusion_matrix(y_test, pred))
print('-----')
print(classification_report(y_test, pred))

WITH K=10
-----
[[2898  103]
 [ 116 481]]
-----
              precision    recall  f1-score   support

      0       0.96       0.97       0.96       3001
      1       0.82       0.81       0.81       597

 accuracy         0.94
 macro avg       0.89
weighted avg       0.94
```

As seen above we are getting the best results with default K (K=5)

### 8.3.5 KNeighbors Classifier GridsearchCV for Choosing Reasonable K Values

In [136...

```
k_values = range(1, 30)
param_grid = {"n_neighbors": k_values, "p": [1, 2], "weights": ['uniform', "distance"]}
```

In [137...

```
KNN_grid = KNeighborsClassifier()
KNN_grid_model = GridSearchCV(KNN_grid, param_grid, cv = 10, scoring = 'recall')
KNN_grid_model.fit(X_train, y_train)
```

Out[137]:

▸ **GridSearchCV**

▸ estimator: KNeighborsClassifier

▸ KNeighborsClassifier

In [138... KNN\_grid\_model.best\_estimator\_

Out[138]: 

KNeighborsClassifier

KNeighborsClassifier(n\_neighbors=3, p=1, weights='distance')

In [139... print(colored('\033[1mBest Parameters of GridSearchCV for KNN Model:\033[0m', 'blue'), colored(KNN\_grid\_model.best\_params\_, 'red'))

Best Parameters of GridSearchCV for KNN Model: {'n\_neighbors': 3, 'p': 1, 'weights': 'distance'}

In [140... 

```
# NOW WITH K=3

KNN_tuned3 = KNeighborsClassifier(n_neighbors = 3, p = 1, weights = 'distance')
KNN_tuned3.fit(X_train, y_train)
y_pred = KNN_tuned3.predict(X_test)
y_train_pred = KNN_tuned3.predict(X_train)

KNN_tuned3_f1 = f1_score(y_test, y_pred)
KNN_tuned3_acc = accuracy_score(y_test, y_pred)
KNN_tuned3_recall = recall_score(y_test, y_pred)
KNN_tuned3_auc = roc_auc_score(y_test, y_pred)

print("KNN_tuned (K=3)")
print ("-----")
eval(KNN_tuned3, X_train, X_test)
train_val(y_train, y_train_pred, y_test, y_pred)
```

```
KNN_tuned (K=3)
-----
[[2897  104]
 [  79  518]]
Test_Set
```

	precision	recall	f1-score	support
0	0.97	0.97	0.97	3001
1	0.83	0.87	0.85	597
accuracy			0.95	3598
macro avg	0.90	0.92	0.91	3598
weighted avg	0.95	0.95	0.95	3598

```
Train_Set
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	6999
1	1.00	1.00	1.00	1394
accuracy			1.00	8393
macro avg	1.00	1.00	1.00	8393
weighted avg	1.00	1.00	1.00	8393

Out[140]:

	train_set	test_set
Accuracy	1.000	0.949
Precision	1.000	0.833
Recall	1.000	0.868
f1	1.000	0.850

In [141... 

```
# NOW WITH K=1

KNN_tuned1 = KNeighborsClassifier(n_neighbors = 1, p = 1, weights = 'distance')
KNN_tuned1.fit(X_train, y_train)
y_pred = KNN_tuned1.predict(X_test)
y_train_pred = KNN_tuned1.predict(X_train)

KNN_tuned1_f1 = f1_score(y_test, y_pred)
KNN_tuned1_acc = accuracy_score(y_test, y_pred)
KNN_tuned1_recall = recall_score(y_test, y_pred)
KNN_tuned1_auc = roc_auc_score(y_test, y_pred)

print("KNN_tuned (K=1)")
print ("-----")
eval(KNN_tuned1, X_train, X_test)
```

```
KNN_tuned (K=1)
-----
[[2882  119]
 [  84  513]]
Test_Set
```

	precision	recall	f1-score	support
0	0.97	0.96	0.97	3001
1	0.81	0.86	0.83	597
accuracy			0.94	3598
macro avg	0.89	0.91	0.90	3598
weighted avg	0.95	0.94	0.94	3598

```
Train_Set
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	6999
1	1.00	1.00	1.00	1394
accuracy			1.00	8393
macro avg	1.00	1.00	1.00	8393
weighted avg	1.00	1.00	1.00	8393

K=1 and K=3 have the same error rate. In order to reduce the complexity of the model, we can be continue with K=1 as the tuned\_model. However, I will continue with tuned\_model K=3 since recall values are better for K=3.

### 8.3.6 Prediction

In [142... 

```
cprint('KNN_tuned Predictions','green', 'on_red')
KNN_Pred = {"Actual": y_test, "KNN_Pred":y_pred}
KNN_Pred = pd.DataFrame.from_dict(KNN_Pred)
KNN_Pred.head()
```

KNN\_tuned Predictions

Out[142]:

	Actual	KNN_Pred
3118	0	0
10490	0	0
1106	1	1
3822	0	0
6873	0	0

In [143... 

```
cprint('Predictions','green', 'on_red')
KNN_Pred.drop("Actual", axis = 1, inplace = True)
Model_Preds = pd.merge(Model_Preds, KNN_Pred, left_index = True, right_index = True)
Model_Preds.head()
```

Predictions

Out[143]:

	Actual	GB_Pred	KNN_Pred
3118	0	0	0
10490	0	0	0
1106	1	1	1
3822	0	0	0
6873	0	0	0

## 8.4 - Random Forest Classifier

### 8.4.1 Model Building

In [144...

```
RF_model = RandomForestClassifier(class_weight = "balanced", random_state = 101)
RF_model.fit(X_train, y_train)
y_pred = RF_model.predict(X_test)
y_train_pred = RF_model.predict(X_train)

RF_model_f1 = f1_score(y_test, y_pred)
RF_model_acc = accuracy_score(y_test, y_pred)
RF_model_recall = recall_score(y_test, y_pred)
RF_model_auc = roc_auc_score(y_test, y_pred)
```

### 8.4.2 Evaluating Model Performance

In [145...

```
print("RF_Model")
print ("-----")
eval(RF_model, X_train, X_test)
```

RF\_Model

-----

[[2996 5]

[ 56 541]]

Test\_Set

	precision	recall	f1-score	support
0	0.98	1.00	0.99	3001
1	0.99	0.91	0.95	597
accuracy			0.98	3598
macro avg	0.99	0.95	0.97	3598
weighted avg	0.98	0.98	0.98	3598

Train\_Set

	precision	recall	f1-score	support
0	1.00	1.00	1.00	6999
1	1.00	1.00	1.00	1394
accuracy			1.00	8393
macro avg	1.00	1.00	1.00	8393
weighted avg	1.00	1.00	1.00	8393

In [146...

```
cprint('RF_model Scores','green', 'on_red')
train_val(y_train, y_train_pred, y_test, y_pred)
```


RF\_model Scores

Out[146]:

	train_set	test_set
Accuracy	1.000	0.983
Precision	1.000	0.991
Recall	1.000	0.906
f1	1.000	0.947

In [147...

```
from yellowbrick.classifier import ClassPredictionError
visualizer = ClassPredictionError(RF_model)
# Fit the training data to the visualizer
visualizer.fit(X_train, y_train)
# Evaluate the model on the test data
visualizer.score(X_test, y_test)
# Draw visualization
visualizer.poof();
```



### 8.4.3 Feature Importance for Random Forest Model

In [148...

```
RF_feature_imp = pd.DataFrame(index=X.columns, data = RF_model.feature_importances_, columns = ['Importance']).sort_values("Importance", ascending = True)
RF_feature_imp
```




Out[148]:

	Importance
satisfaction_level	0.282
time_spend_company	0.233
average_montly_hours	0.155
last_evaluation	0.147
number_project	0.136
work_accident	0.012
salary_low	0.007
salary_medium	0.005
department_sales	0.004
department_technical	0.004
department_support	0.004
department_hr	0.002
department_RandD	0.002
department_accounting	0.002
department_management	0.002
department_marketing	0.001
promotion_last_5years	0.001
department_product_mng	0.001

In [149...

fig = px.bar(RF\_feature\_imp.sort\_values('Importance', ascending = False), x = RF\_feature\_imp.sort\_values('Importance', ascending = False).index, y = 'Importance', title = "Feature Importance", labels = dict(x = "Features", y ="Feature\_Importance"))  
fig.show()



### 8.4.4 Random Forest Classifier Cross Validation

In [150...

RF\_cv = RandomForestClassifier(class\_weight = "balanced", random\_state = 101)  
RF\_cv\_scores = cross\_validate(RF\_cv, X\_train, y\_train,  
 scoring = ['accuracy', 'precision', 'recall', 'f1', 'roc\_auc'], cv = 10)  
RF\_cv\_scores = pd.DataFrame(RF\_cv\_scores, index = range(1, 11))  
RF\_cv\_scores.mean()[2:]

Out[150]:

test\_accuracy 0.982  
test\_precision 0.982  
test\_recall 0.906  
test\_f1 0.942  
test\_roc\_auc 0.981  
dtype: float64

### 8.4.5 Random Forest Classifier GridSearchCV

In [151...

param\_grid = {'n\_estimators' : [50, 100, 300],  
 'max\_features' : [2, 3, 4],  
 'max\_depth' : [3, 5, 7, 9],  
 'min\_samples\_split' : [2, 5, 8]}

In [152...

RF\_grid = RandomForestClassifier(class\_weight = 'balanced', random\_state = 101)  
RF\_grid\_model = GridSearchCV(estimator = RF\_grid,  
 param\_grid = param\_grid,  
 scoring = "recall",  
 n\_jobs = -1, verbose = 2)  
RF\_grid\_model.fit(X\_train, y\_train)  
  
Fitting 5 folds for each of 108 candidates, totalling 540 fits

Out[152]:

GridSearchCV  
estimator: RandomForestClassifier  
RandomForestClassifier

In [153...

RF\_grid\_model.best\_estimator\_

Out[153]:

RandomForestClassifier  
RandomForestClassifier(class\_weight='balanced', max\_depth=3, max\_features=4,  
 n\_estimators=300, random\_state=101)

In [154...

print(colored('\033[1mBest Parameters of GridSearchCV for Random Forest Model:\033[0m', 'blue'), colored(RF\_grid\_model.best\_params\_, 'red'))  
  
Best Parameters of GridSearchCV for Random Forest Model: {'max\_depth': 3, 'max\_features': 4, 'min\_samples\_split': 2, 'n\_estimators': 300}

In [155...

RF\_tuned = RandomForestClassifier(class\_weight = 'balanced',  
 max\_depth = 3,  
 max\_features = 4,  
 n\_estimators = 300,  
 min\_samples\_split = 2,  
 random\_state = 101).fit(X\_train, y\_train)



```
In [156... y_pred = RF_tuned.predict(X_test)
y_train_pred = RF_tuned.predict(X_train)

RF_tuned_f1 = f1_score(y_test, y_pred)
RF_tuned_acc = accuracy_score(y_test, y_pred)
RF_tuned_recall = recall_score(y_test, y_pred)
RF_tuned_auc = roc_auc_score(y_test, y_pred)
```

```
In [157... print("RF_tuned")
print ("-----")
eval(RF_tuned, X_train, X_test)

RF_tuned
-----
[[2804  197]
 [  47 550]]
Test_Set
```

	precision	recall	f1-score	support
0	0.98	0.93	0.96	3001
1	0.74	0.92	0.82	597
accuracy			0.93	3598
macro avg	0.86	0.93	0.89	3598
weighted avg	0.94	0.93	0.94	3598

```
Train_Set
```


	precision	recall	f1-score	support
0	0.99	0.93	0.96	6999
1	0.73	0.94	0.82	1394
accuracy			0.93	8393
macro avg	0.86	0.93	0.89	8393
weighted avg	0.94	0.93	0.94	8393

```
In [158... cprint('RF_tuned Scores','green', 'on_red')
train_val(y_train, y_train_pred, y_test, y_pred)
```

RF\_tuned Scores

	train_set	test_set
Accuracy	0.932	0.932
Precision	0.731	0.736
Recall	0.937	0.921
f1	0.821	0.818

```
In [159... from yellowbrick.classifier import ClassPredictionError
visualizer = ClassPredictionError(RF_tuned)
# Fit the training data to the visualizer
visualizer.fit(X_train, y_train)
# Evaluate the model on the test data
visualizer.score(X_test, y_test)
# Draw visualization
visualizer.poof();
```



8.4.6 Prediction

```
In [160... cprint('RF_tuned Predictions','green', 'on_red')
RF_Pred = {"Actual": y_test, "RF_Pred":y_pred}
RF_Pred = pd.DataFrame.from_dict(RF_Pred)
RF_Pred.head()
```

RF\_tuned Predictions

	Actual	RF_Pred
3118	0	0
10490	0	0
1106	1	1
3822	0	0
6873	0	0

```
In [161... cprint('Predictions','green', 'on_red')
RF_Pred.drop("Actual", axis = 1, inplace = True)
Model_Preds = pd.merge(Model_Preds, RF_Pred, left_index = True, right_index = True)
Model_Preds.head()
```

Predictions

	Actual	GB_Pred	KNN_Pred	RF_Pred
3118	0	0	0	0
10490	0	0	0	0
1106	1	1	1	1
3822	0	0	0	0
6873	0	0	0	0

## 8.5 - CatBoost Classifier

### 8.5.1 Model Building

```
In [162... CB_model = CatBoostClassifier(verbose = False, scale_pos_weight = 4, random_state = 101)
CB_model.fit(X_train, y_train)
y_pred = CB_model.predict(X_test)
y_train_pred = CB_model.predict(X_train)

CB_model_f1 = f1_score(y_test, y_pred)
CB_model_acc = accuracy_score(y_test, y_pred)
CB_model_recall = recall_score(y_test, y_pred)
CB_model_auc = roc_auc_score(y_test, y_pred)
```

### 8.5.2 Evaluating Model Performance

```
In [163... print("CB_Model")
print ("-----")
eval(CB_model, X_train, X_test)
```

```
CB_Model
-----
[[2977  24]
 [ 47 550]]
Test_Set
```

	precision	recall	f1-score	support
0	0.98	0.99	0.99	3001
1	0.96	0.92	0.94	597
accuracy			0.98	3598
macro avg	0.97	0.96	0.96	3598
weighted avg	0.98	0.98	0.98	3598

```
Train_Set
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	6999
1	0.98	0.99	0.98	1394
accuracy			0.99	8393
macro avg	0.99	0.99	0.99	8393
weighted avg	0.99	0.99	0.99	8393

```
In [164... cprint('CB_model Scores','green', 'on_red')
train_val(y_train, y_train_pred, y_test, y_pred)
```

CB\_model Scores

```
Out[164]:
```


	train_set	test_set
Accuracy	0.994	0.980
Precision	0.977	0.958
Recall	0.987	0.921
f1	0.982	0.939

```
In [165... CB_cm = confusion_matrix(y_test, y_pred)
CB_cm_df = pd.DataFrame(CB_cm)
CB_cm_df = CB_cm_df.rename(columns={0:"Employee_Stayed", 1:"Employee_Left"}, index={0:"Employee_Stayed", 1:"Employee_Left"})
CB_cm_df["Total"] = CB_cm_df["Employee_Stayed"] + CB_cm_df["Employee_Left"]
```

```
In [166... fig = px.bar(CB_cm_df, x="Employee_Stayed", y="Total", color="Employee_Left", title="CatBoost Confusion Matrix Distribution")

fig.update_layout(
    xaxis_title="Employees-Left",
    yaxis_title="The Number of Employees",
    font=dict(
        family="Courier New, monospace",
        size=14,
        color="#7f7f7f"
    )
)

fig.show()
```




### 8.5.3 Feature Importance for CatBoost Model

```
In [167... CB_feature_imp = pd.DataFrame(index = X.columns, data = CB_model.feature_importances_, columns = ['Importance']).sort_values("Importance", ascending=True)
CB_feature_imp
```

```
Out[167]:
```

	Importance
satisfaction_level	26.423
time_spend_company	20.504
number_project	15.951
average_monthly_hours	14.938
last_evaluation	13.180
salary_low	1.730
work_accident	1.684
salary_medium	1.175
department_support	0.892
department_sales	0.886
department_product_mng	0.618
department_hr	0.528
department_technical	0.466
department_accounting	0.370
department_RandD	0.260
department_management	0.181
promotion_last_5years	0.124
department_marketing	0.091

```
In [168... fig = px.bar(CB_feature_imp.sort_values('Importance', ascending = False), x = CB_feature_imp.sort_values('Importance', ascending = False).index, y = 'Importance', title = "Feature Importance", labels = dict(x = "Features", y ="Feature_Importance"))
fig.show()
```



8.5.4 CatBoost Classifier Cross Validation

```
In [169... CB_cv = CatBoostClassifier(verbose = False, scale_pos_weight = 4, random_state = 101)
CB_cv_scores = cross_validate(CB_cv, X_train, y_train,
                              scoring = ['accuracy', 'precision', 'recall', 'f1', 'roc_auc'], cv = 10)
CB_cv_scores = pd.DataFrame(CB_cv_scores, index = range(1, 11))

CB_cv_scores.mean()[2:]

Out[169]: test_accuracy    0.980
test_precision    0.952
test_recall       0.924
test_f1           0.938
test_roc_auc      0.983
dtype: float64
```

8.5.5 CatBoost Classifier GridSearchCV

```
In [170... param_grid = {'learning_rate': [0.01, 0.03, 0.1, 0.5],
                'depth': [4, 6, 8, 10],
                'l2_leaf_reg': [1, 3, 5, 7, 9]}

In [171... CB_grid = CatBoostClassifier(verbose = False, scale_pos_weight = 4, random_state = 101)
CB_grid_model = GridSearchCV(estimator = CB_grid,
                             param_grid = param_grid,
                             scoring = "recall",
                             n_jobs = -1, verbose = 2)
CB_grid_model.fit(X_train, y_train)

Fitting 5 folds for each of 80 candidates, totalling 400 fits

Out[171]:  GridSearchCV
 estimator: CatBoostClassifier
           CatBoostClassifier

In [172... CB_grid_model.best_params_

Out[172]: {'depth': 4, 'l2_leaf_reg': 3, 'learning_rate': 0.01}

In [173... print(colored('\033[1mBest Parameters of GridSearchCV forCatBoost Model:\033[0m', 'blue'), colored(CB_grid_model.best_params_, 'red'))

Best Parameters of GridSearchCV forCatBoost Model: {'depth': 4, 'l2_leaf_reg': 3, 'learning_rate': 0.01}

In [174... CB_tuned = CatBoostClassifier(verbose = False,
                                scale_pos_weight = 4,
                                depth = 4,
                                l2_leaf_reg = 3,
                                learning_rate = 0.01,
                                random_state = 101).fit(X_train, y_train)
```

```
In [175... y_pred = CB_tuned.predict(X_test)
y_train_pred = CB_tuned.predict(X_train)

CB_tuned_f1 = f1_score(y_test, y_pred)
CB_tuned_acc = accuracy_score(y_test, y_pred)
CB_tuned_recall = recall_score(y_test, y_pred)
CB_tuned_auc = roc_auc_score(y_test, y_pred)

In [176... print("CB_tuned")
print ("-----")
eval(CB_tuned, X_train, X_test)

CB_tuned
-----
[[2946  55]
 [  45 552]]
Test_Set
      precision    recall  f1-score   support

      0       0.98       0.98       0.98        3001
      1       0.91       0.92       0.92         597

   accuracy          0.97
  macro avg       0.95       0.95       0.95
weighted avg       0.97       0.97       0.97

Train_Set
      precision    recall  f1-score   support

      0       0.99       0.98       0.99        6999
      1       0.92       0.94       0.93       1394

   accuracy          0.98
  macro avg       0.96       0.96       0.96
weighted avg       0.98       0.98       0.98

In [177... cprint('CB_tuned Scores','green', 'on_red')
train_val(y_train, y_train_pred, y_test, y_pred)

CB_tuned Scores
Out[177]:


|                  | train_set | test_set |
|------------------|-----------|----------|
| <b>Accuracy</b>  | 0.977     | 0.972    |
| <b>Precision</b> | 0.923     | 0.909    |
| <b>Recall</b>    | 0.940     | 0.925    |
| <b>f1</b>        | 0.931     | 0.917    |




In [178... CB_cm = confusion_matrix(y_test, y_pred)
CB_cm_df = pd.DataFrame(CB_cm)
CB_cm_df = CB_cm_df.rename(columns={0:"Employee_Stayed", 1:"Employee_Left"}, index={0:"Employee_Stayed", 1:"Employee_Left"})
CB_cm_df["Total"] = CB_cm_df["Employee_Stayed"] + CB_cm_df["Employee_Left"]

In [179... fig = px.bar(CB_cm_df, x="Employee_Stayed", y="Total", color="Employee_Left", title="CatBoost Confusion Matrix Distribution")

fig.update_layout(
    xaxis_title="Employees-Left", Employees_Stayed",
    yaxis_title="The Number of Employees",
    font=dict(
        family="Courier New, monospace",
        size=14,
        color="#7f7f7f"
    )
)

fig.show()
```



8.5.6 Prediction

```
In [180... cprint('CB_tuned Predictions','green', 'on_red')
CB_Pred = {"Actual": y_test, "CB_Pred":y_pred}
CB_Pred = pd.DataFrame.from_dict(CB_Pred)
CB_Pred.head()

CB_tuned Predictions
Out[180]:


|              | Actual | CB_Pred |
|--------------|--------|---------|
| <b>3118</b>  | 0      | 0       |
| <b>10490</b> | 0      | 0       |
| <b>1106</b>  | 1      | 1       |
| <b>3822</b>  | 0      | 0       |
| <b>6873</b>  | 0      | 0       |



In [181... cprint('Predictions','green', 'on_red')
CB_Pred.drop("Actual", axis = 1, inplace = True)
Model_Preds = pd.merge(Model_Preds, CB_Pred, left_index = True, right_index = True)
Model_Preds.head()

Predictions
```

Out[181]:

	Actual	GB_Pred	KNN_Pred	RF_Pred	CB_Pred
3118	0	0	0	0	0
10490	0	0	0	0	0
1106	1	1	1	1	1
3822	0	0	0	0	0
6873	0	0	0	0	0

In [182...]

cprint('Random Predictions','green', 'on\_red')  
Model\_Preds.sample(10)

Random Predictions


Out[182]:




	Actual	GB_Pred	KNN_Pred	RF_Pred	CB_Pred
1253	1	1	1	1	1
239	1	1	1	1	1
10283	0	0	0	0	0
10044	0	0	0	0	0
11144	0	0	0	0	0
779	1	1	1	1	1
5321	0	0	0	0	0
2017	0	0	0	0	0
7664	0	0	1	0	0
10710	0	0	0	0	0

## 9 - THE COMPARISON OF MODELS

In [183...

compare = pd.DataFrame({"Model": ["GB\_model", "GB\_tuned", "KNN\_Model", "KNN\_tuned3", "RF\_model", "RF\_tuned", "CB\_model",  
"CB\_tuned"],  
  
"F1\_Score": [GB\_model\_f1, GB\_tuned\_f1, KNN\_model\_f1, KNN\_tuned3\_f1, RF\_model\_f1, RF\_tuned\_f1,  
CB\_model\_f1, CB\_tuned\_f1],  
  
"Accuracy\_Score": [GB\_model\_acc, GB\_tuned\_acc, KNN\_model\_acc, KNN\_tuned3\_acc, RF\_model\_acc,  
RF\_tuned\_acc, CB\_model\_acc, CB\_tuned\_acc],  
  
"Recall\_Score": [GB\_model\_recall, GB\_tuned\_recall, KNN\_model\_recall, KNN\_tuned3\_recall, RF\_model\_recall,  
RF\_tuned\_recall, CB\_model\_recall, CB\_tuned\_recall],  
  
"ROC\_AUC\_Score": [GB\_model\_auc, GB\_tuned\_auc, KNN\_model\_auc, KNN\_tuned3\_auc, RF\_model\_auc,  
RF\_tuned\_auc, CB\_model\_auc, CB\_tuned\_auc]})  
  
compare = compare.sort\_values(by="Recall\_Score", ascending=True)  
fig = px.bar(compare, x = "Recall\_Score", y = "Model", title = "Recall\_Score")  
fig.show()  
  
compare = compare.sort\_values(by="F1\_Score", ascending=True)  
fig = px.bar(compare, x = "F1\_Score", y = "Model", title = "F1\_Score")  
fig.show()  
  
compare = compare.sort\_values(by="Accuracy\_Score", ascending=True)  
fig = px.bar(compare, x = "Accuracy\_Score", y = "Model", title = "Accuracy\_Score")  
fig.show()  
  
compare = compare.sort\_values(by="ROC\_AUC\_Score", ascending=True)  
fig = px.bar(compare, x = "ROC\_AUC\_Score", y = "Model", title = "ROC\_AUC\_Score")  
fig.show()





In [184...

```
cprint('Scores', 'green', 'on_red')
compare.T
```

Out[184]:

Scores									
	2	3	5	4	7	6	0	1	
Model	KNN_Model	KNN_tuned3	RF_tuned	RF_model	CB_tuned	CB_model	GB_model	GB_tuned	
F1_Score	0.836	0.850	0.818	0.947	0.917	0.939	0.939	0.952	
Accuracy_Score	0.944	0.949	0.932	0.983	0.972	0.980	0.980	0.985	
Recall_Score	0.861	0.868	0.921	0.906	0.925	0.921	0.925	0.920	
ROC_AUC_Score	0.911	0.917	0.928	0.952	0.953	0.957	0.958	0.959	

# 10 - MODEL DEPLOYMENT

- Save and Export the Model as .pkl
- Save and Export Variables as .pkl

You cooked the food in the kitchen and moved on to the serving stage. The question is how do you showcase your work to others? Model Deployment helps you showcase your work to the world and make better decisions with it. But, deploying a model can get a little tricky at times. Before deploying the model, many things such as data storage, preprocessing, model building and monitoring need to be studied. Streamlit is a popular open source framework used by data scientists for model distribution.

Deployment of machine learning models, means making your models available to your other business systems. By deploying models, other systems can send data to them and get their predictions, which are in turn populated back into the company systems. Through machine learning model deployment, can begin to take full advantage of the model you built.

Data science is concerned with how to build machine learning models, which algorithm is more predictive, how to design features, and what variables to use to make the models more accurate. However, how these models are actually used is often neglected. And yet this is the most important step in the machine learning pipeline. Only when a model is fully integrated with the business systems, real values can be extract from its predictions.

After doing the following operations in this notebook, jump to new .py file and create your web app with Streamlit.

In [185...

gradient\_boosting\_classifier = pickle.dump(GB\_tuned, open('gradient\_boosting\_model', 'wb'))

In [186...

kneighbors\_classifier = pickle.dump(KNN\_tuned3, open('kneighbors\_model', 'wb'))

In [187...

random\_forest\_classifier = pickle.dump(RF\_tuned, open('random\_forest\_model', 'wb'))

In [188...

catboost\_classifier = pickle.dump(CB\_tuned, open('catboost\_model', 'wb'))

## 10.2 - Save and Export Variables as .pkl

In [189...

```
# col_lst = ['satisfaction_level', 'last_evaluation', 'number_project', 'average_monthly_hours', 'time_spend_company',
#           'work_accident', 'promotion_last_5years', 'department_RandD', 'department_accounting', 'department_hr',
#           'department_management', 'department_marketing', 'department_product_mng', 'department_sales',
#           'department_support', 'department_technical', 'salary_Low', 'salary_medium']
# scaler = MinMaxScaler()
# scaler_fitted = scaler.fit(df2[col_lst])
# scaler_deploy = pickle.dump(scaler_fitted, open('scaler.sav', 'wb'))
```

# 11 - CONCLUSION

In this project we have HR data of a company. A study is requested from us to predict which employee will churn by using this data.

First of all, to observe the structure of the data, outliers, missing values and features that affect the target variable, we used exploratory data analysis and data visualization techniques.

Then, we performed data pre-processing operations such as **\*Scaling\*** and **\*Label Encoding\*** to increase the accuracy score of Gradient Descent Based or Distance-Based algorithms.\*\*

We used the **\*K-means\*** algorithm to make cluster analysis. In order to find the optimal number of clusters, we used the **\*Elbow method\***. Briefly, tried to predict the set to which individuals were related by using K-means and evaluate the estimation results.

Then we built models to predict whether employees will churn or not. We trained our models with train set, tested the success of models with test set.

In this study, we made modelling with **\*Gradient Boosting Classifier, K Neighbors Classifier, Random Forest Classifier\*** and **\*CatBoost Classifier\***.

We used scikit-learn **\*Confusion Metrics\*** module for accuracy calculation and the **\*Yellowbrick\*** module for model selection and visualization.

In [ ]:

In [ ]: